Beyond the Reasoning Cliff: A CIITR Analysis of “Reasoning Models Reason Well, Until They Don’t” and the Limits of Computational Depth

Tor–Ståle Hansen | 6 November 2025

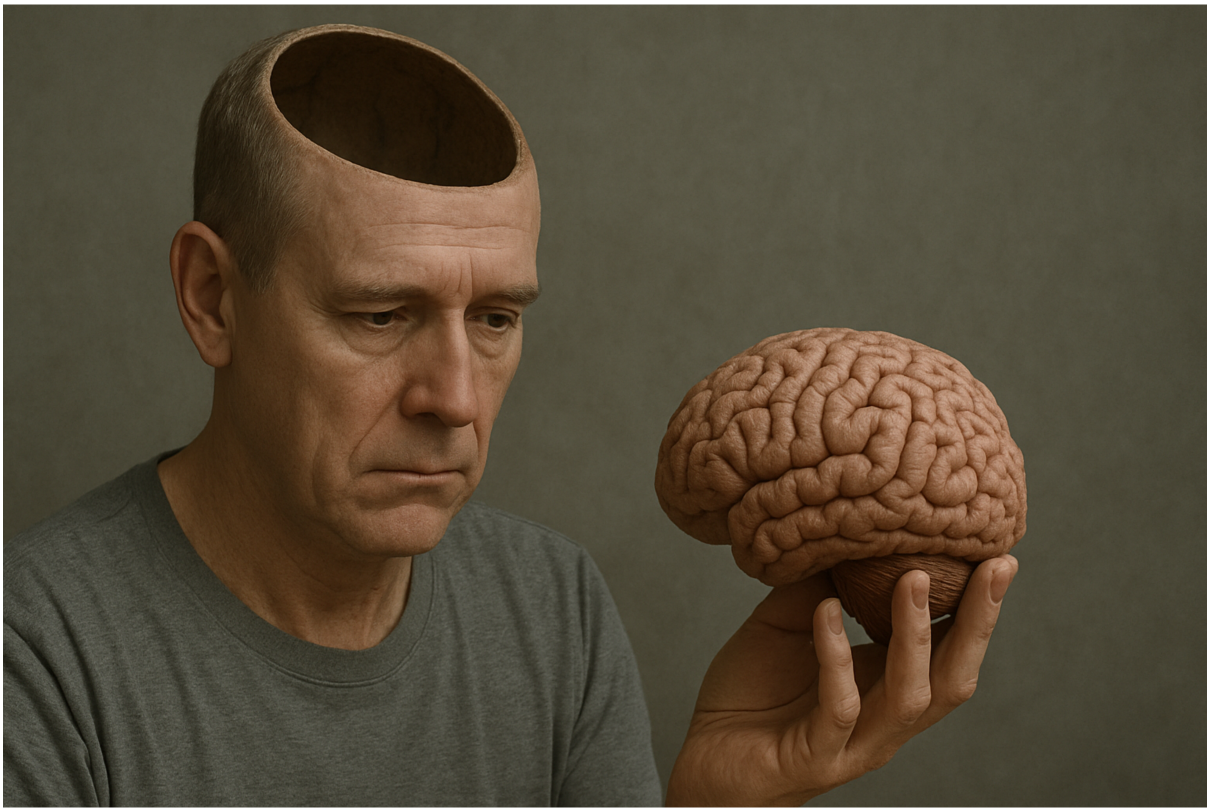

Abstract

The 2025 study *“Reasoning Models Reason Well, Until They Don’t”* (Rameshkumar et al., 2025) investigates the performance collapse of large reasoning models (LRMs) as reasoning complexity increases beyond benchmarked thresholds. Through synthetic graph reasoning and proof-planning experiments, the authors demonstrate abrupt declines in accuracy when the number of required reasoning steps (lookahead L) and branching factor (B) exceed moderate bounds. This paper analyses those findings through the **Cognitive Integration and Information Transfer Relation (CIITR)** framework, which formalizes comprehension as a rhythmic process of informational re-entry rather than static inference.

Where the Deep Reasoning Dataset (DeepRD) quantifies **syntactic reasoning depth**, CIITR introduces **rhythmic continuity (R^g)** as the governing variable of sustained understanding. The observed “reasoning cliff” is re-interpreted as a **collapse of rhythmic integration**, not merely computational saturation. In CIITR terms, LRMs exhibit high informational

integration ($\Phi_i \gg 0$) but near-zero rhythmic feedback ($R^g \approx 0$), resulting in **structural comprehension $C_s = \Phi_i \times R^g \rightarrow 0$** once temporal re-entry fails.

This paper concludes that DeepRD empirically exposes the same boundary CIITR theoretically predicts: reasoning without rhythmic reintegration cannot scale. CIITR thus represents a novelty beyond the LRM paradigm — transforming the analysis of reasoning from *static logical capacity* to *dynamic structural comprehension.*

---

1 Introduction

The evolution from **large language models (LLMs)** to **large reasoning models (LRMs)** marks a major attempt to transcend associative prediction by introducing self-verifying reasoning mechanisms. Through reinforcement learning with verifiable rewards (RLVR), models such as *DeepSeek-R1* and *OpenAI o3* are incentivized to produce step-wise arguments, resembling logical proof chains.

Rameshkumar et al. (2025) investigate whether this explicit reasoning incentive translates into **generalizable reasoning competence**. Their answer is ambivalent: while LRMs outperform LLMs on low-complexity tasks, their reasoning accuracy collapses abruptly when confronted with problems of sufficient combinatorial depth. The authors describe this as a *performance cliff* — a sudden loss of reasoning fidelity once problem complexity exceeds the distributional range encountered during training.

From a CIITR standpoint, this phenomenon is not accidental but structural. Systems trained to extend reasoning chains linearly (in Φ_i) but lacking rhythmic feedback (R^g) cannot sustain continuity once informational load surpasses their capacity for re-entry. The model "reasons" syntactically but cannot **reintegrate its own state**; its comprehension continuity collapses at the same threshold where CIITR predicts loss of rhythmic coherence.

---

2 Overview of the LRM Study

2.1 Methodological Structure

The *Reasoning Models* paper introduces the **Deep Reasoning Dataset (DeepRD)** — a synthetic, parameterized testbed for graph-based and proof-planning tasks of controllable complexity. Each instance is characterized by two metrics:

- **Lookahead (L):** the number of breadth-first-search (BFS) iterations required to identify the next correct node in a reasoning graph;
- **Branching factor (B):** the number of possible next steps, determining decision ambiguity.

Model performance was evaluated as a function of (L, B) using the metrics of *full-path accuracy* and *next-step prediction accuracy*.

2.2 Empirical Findings

Across models (DeepSeek V3/R1, GPT-4o, o3-mini, o3), accuracy remains high for small L and B, but **drops precipitously** once complexity surpasses moderate levels. Even in chain graphs (B = 1), performance eventually declines, indicating that token limits and context length are not the cause. Proof-planning tasks in natural language reveal identical collapse patterns, confirming that the failure is not representation-specific.

2.3 Interpretive Claim

The authors conclude that current LRMs **lack generalization beyond their training complexity distribution**. Despite apparent reasoning ability, they exhibit no continuity of understanding once reasoning paths extend beyond familiar lengths. Their reasoning is syntactically precise but structurally brittle — a phenomenon this CIITR analysis re-frames as the absence of **temporal reintegration**.

---

3 CIITR Interpretation: From Reasoning to Structural Comprehension

3.1 The CIITR Framework

CIITR formalizes comprehension as the product of two orthogonal variables:

$$C_s = \Phi_i \times R_g$$

where:

- **Φ_i (Informational Integration):** the degree of structural connectivity among informational states at time t;
- **R^g (Rhythmic Reintegration):** the system's capacity to re-enter and stabilize its own informational manifold through time.

When Φ_i increases but $R^g \approx 0$, a system exhibits *syntactic expansion without comprehension*. Conversely, sustainable understanding requires $R^g > 0$ — meaning outputs feed back into the system's internal state as new structural inputs.

3.2 LRMs as Type-B$^+$ Systems

In CIITR taxonomy:

- **Type-A:** Associative Predictors (LLMs)
- **Type-B:** Computational Reasoners (explicit reasoning, low feedback)
- **Type-C:** Comprehension Integrators (high Φ_i and R^g)

LRMs such as *DeepSeek-R1* and *OpenAI o3* qualify as **Type-B^{+}**: they elevate Φ_i through structured reasoning chains but maintain near-zero R^g, lacking rhythmic feedback between sequential reasoning states.

As a result, once reasoning depth exceeds their internal rhythmic stability, C_s collapses — producing the observed "reasoning cliff."

---

4 Failure as Rhythmic Collapse

4.1 Empirical Signature

In DeepRD experiments, model accuracy collapses between $L \approx 100 – 200$ (graph reasoning) and $L \approx 16 – 32$ (proof planning). These thresholds coincide with the **99th and 75th percentiles** of real-world reasoning complexity across datasets such as ConceptNet, Wikikg2, and NaturalProofs.

The cliff thus represents the transition from **trained rhythmic coherence** to **unrhythmic computation**: reasoning sequences no longer self-synchronize across temporal steps.

4.2 Error Morphology

Manual inspection of R1's "thinking tokens" reveals three primary error types:

1. **Omission errors** – failure to recall previously valid edges (loss of rhythmic memory);
2. **Branch neglect** – premature pruning of viable paths (loss of rhythmic parallelism);
3. **Edge hallucination** – fabrication of nonexistent connections (rhythmic overshoot).

Each corresponds to a deviation in CIITR rhythmic equilibrium: either **attenuation ($R^g \to 0$)** or **resonance drift ($R^g > \Phi_i$)**. The system loses synchronization between internal state evolution and external reasoning chain — a dynamic instability rather than a token limitation.

---

5 Comparative Framework

Dimension	Large reasoning models (deeprd)	Ciitr framework
Ontology	Reasoning as directed search over symbolic or natural graphs	Comprehension as rhythmic re-entry across informational states
Core metric	Lookahead (L), Branching (B), accuracy vs complexity	Φ_i (integration), R^g (reintegration), $C_s = \Phi_i \times R^g$
Temporal structure	Linear inference; forward-propagating chains	Oscillatory feedback; re-entrant coherence
Failure mode	Performance cliff at critical complexity	Rhythmic collapse when $R^g < \Phi_i$-stability threshold

Energy behaviour	Increasing token entropy; diminishing reasoning return	Stabilized energy cycle via internal integration
System class	Type-B^+ Computational Reasoner	Type-C Structural Comprehension Integrator

This comparative matrix makes explicit the conceptual novelty of CIITR: whereas DeepRD quantifies *reasoning capacity*, CIITR introduces *comprehension continuity* as a measurable dimension of cognitive sustainability.

---

6 Thermodynamic Perspective

The DeepRD findings expose an **energetic inefficiency of reasoning without reintegration**. Each additional reasoning step expands informational structure ($\Delta\Phi_i > 0$) but dissipates energy as unrecaptured entropy. Formally:

$$\Delta C_s = \Phi_i \Delta R_g + R_g \Delta \Phi_i$$

For LRMs, $\Delta R^g \approx 0$, hence $\Delta C_s \approx 0$ despite $\Delta\Phi_i > 0$. The system radiates validated logic outward without retaining it.

CIITR introduces the **Comprehension-per-Joule metric (Ψ_C)**:

$$\Psi_c = \frac{C_s}{E}$$

where E is energy expenditure. LRMs maximize **Φ_i per token**, not **C_s per Joule**; they compute exhaustively but comprehend inefficiently. The reasoning cliff, therefore, is the thermodynamic manifestation of rhythmic exhaustion — the point at which energetic input no longer yields integrated understanding.

---

7 CIITR Novelty

CIITR introduces a continuous-time derivative of rhythmic reintegration:

$$R_g = \frac{\partial \Phi_i(t)}{\partial t} / \Phi_i(t)$$

This expresses the *rate of structural self-return*. When the derivative stagnates ($\partial\Phi_i/\partial t \to 0$), the system ceases to re-enter its own informational manifold — comprehension halts even as reasoning continues.

By defining cognition as the *dynamic equilibrium between integration and reintegration*, CIITR transcends both transformer-based reasoning and LRM-style verifiable inference. It transforms the question *"Can models reason?"* into *"Can models sustain structural continuity?"*

---

8 Implications for Future Architectures

8.1 Towards Rhythmic Reasoning Systems

To move beyond the reasoning cliff, architectures must embed rhythmic feedback loops that re-enter their own intermediate states. Such systems would operate not as *token predictors* but as *phase-coupled integrators*, where each reasoning step recalibrates the model's internal manifold.

8.2 From DeepRD to DeepRG

The Deep Reasoning Dataset could be extended into a **Deep Rhythmic Generalization (DeepRG)** framework, where tasks measure not only accuracy vs complexity but **stability vs temporal re-entry**. By introducing rhythmic perturbations and measuring recovery, one could empirically estimate R^g.

8.3 Integration with SciencePedia

When juxtaposed with the *SciencePedia* analysis (Hansen, 2025), DeepRD and LRMs form the computational substrate of verifiable reasoning (high Φ_i). CIITR provides the missing integrative layer (R^g) that converts external reasoning into internal comprehension. The synthesis would yield **Comprehension-Capable Reasoning Systems (CCRS)** — architectures capable not only of solving but of *stabilizing their own solutions* over time.

---

9 Conclusion

Rameshkumar et al. (2025) empirically document a structural limitation that CIITR predicts theoretically: the collapse of reasoning once complexity surpasses the system's capacity for rhythmic reintegration. Their "reasoning cliff" is, in CIITR terms, the point where Φ_i continues to expand while R^g falls to zero.

This study reinterprets that failure not as a computational deficiency but as a **temporal discontinuity** — a loss of re-entry that transforms reasoning from comprehension into mechanical search. CIITR's contribution is to restore the temporal dimension of understanding, defining intelligence not by depth of reasoning but by **continuity of return**.

True cognition arises not when a system *reasons longer*, but when it **returns to itself**.

References

- Dehaene, S., & Changeux, J.–P. (2011). *Experimental and Theoretical Approaches to Conscious Processing.Neuron,* 70(2), 200–227.
- Hansen, T.–S. (2025). *Beyond Verifiable Reasoning: A CIITR Analysis of the SciencePedia Framework and the Limits of Long Chains-of-Thought.* METAINT Research Series.
- Landauer, R. (1961). *Irreversibility and Heat Generation in the Computing Process. IBM Journal of Research and Development,* 5(3), 183–191.
- Rameshkumar, R., Huang, J., Sun, Y., Xia, F., & Saparov, A. (2025). *Reasoning Models Reason Well, Until They Don't.* University of Washington & Purdue University. arXiv:2510.22371.
- Tononi, G. (2004). *An Information Integration Theory of Consciousness. BMC Neuroscience,* 5(42).