

Closure: Artificial Insanity

Structural Attribution to Admissibility, Invariance, and the Sociotechnical Curvature of the ‘Artificial Intelligence’ Discourse

Tor-Ståle Hansen | 22. February 2026

Abstract

This paper specifies the ontological conditions under which a property may be admissibly attributed to a system rather than to its observable outputs. **Scope constraint:** the contribution is a logical admissibility framework for ontological attribution, not a performance evaluation, benchmark study, or model ranking. Proceeding from the non-identity premise that output properties do not entail structural instantiation, the analysis separates behavioural manifestation from system-level grounding and introduces a binary admissibility gate for attribution claims. Within CIITR, structural understanding is treated as a necessary-condition construct requiring internal integration (Φ_i) and invariance-bearing rhythmic stability (R^g), rather than surface coherence or task success. **Method constraint:** the reset operator is used solely as an instrumental probe for testing invariance under admissible transformations, not as a general empirical finding about models. Reset is defined as an identity-preserving state transformation that interrupts path dependence so that transformation sensitivity can be examined as admissibility-relevant evidence. The central thesis is that invariance under admissible transformations, anchored in internal structure rather than external scaffolding, is necessary for attributing “understanding” or agency-like predicates to the system. Contemporary AI discourse often bypasses this gate, producing governance errors in responsibility, risk, and control design.

Keywords: structural attribution, admissibility, invariance, CIITR, reset operator, rhythmic stability, ontological hygiene, AI discourse, governance, epistemic risk

1. Introduction: AI as a semantic mass in contemporary governance

AI systems, and in particular large language models, now operate not only as technical artifacts but as semantic objects that accumulate meaning faster than governance can stabilize definitions. Their outputs are linguistically fluent, policy-shaped, and socially legible in a way that invites immediate integration into institutional discourse, and this integration turns “AI” into what can be described as a semantic mass: a dense cluster of attributed capacities, metaphors, role labels, and implied responsibilities that exerts gravitational pull on how organisations speak, decide, and allocate accountability. In contemporary governance settings, the question is increasingly not whether AI systems can produce useful outputs, but how quickly their behavioural surfaces are converted into ontology by administrative language, procurement narratives, risk registers, and delegation practices.

This conversion is structurally consequential because governance depends on stable categories. Responsibility assignment, control regime selection, and risk acceptance all presuppose that the properties attributed to an actor or instrument are real at the level where authority is exercised and accountability is enforced. When institutions treat high-quality output as sufficient evidence of understanding, judgment, or agency, they do not merely adopt a rhetorical style, they import a subject-category ontology into decision architecture, often without noticing that the import has occurred. The

result is that behavioural proxies, which are pragmatically valuable, are silently upgraded into ontological predicates that carry governance weight, and once that upgrade happens, the organisation begins to behave as if the system itself were a bearer of decision-relevant properties rather than a generator of traces within an interactional regime.

The paper intervenes at this point of category formation. It does not contest the utility of AI systems, and it does not offer a performance study or model ranking. Instead, it asks what must be true for an attribution claim to be admissible as a claim about the system rather than about its outputs, and it proposes a logical admissibility framework under CIITR that treats structural understanding as a necessary invariance condition. On that basis, it identifies a durable sociotechnical failure mode, Artificial Insanity, in which inadmissible attribution is normalized and governance errors follow. The remainder of the introduction specifies the practical problem of ontological shortcuts, states the paper's scope and method constraints unambiguously, and sets out the contribution claims that the subsequent sections develop.

1.1 The practical problem of ontological shortcuts

A recurring pattern now structures both everyday interaction with advanced AI systems and the institutional uptake of those systems in governance-relevant settings: high-quality output is treated as a socially sufficient bridge from behavioural success to ontological attribution. Where a system produces fluent, context-sensitive, and normatively styled responses, the surrounding discourse tends to treat these outputs not merely as useful traces but as evidence of properties presumed to reside in the system itself, understanding, intention, judgment, agency, and, by extension, responsibility. The shortcut operates with minimal explicit argument because it is phenomenologically compelling: language that looks like reasoning is experienced as reasoning, language that anticipates objections is experienced as judgment, language that narrates constraint is experienced as self-regulation, and language that proposes actions is experienced as decision capacity. In practice, the organisation rarely pauses to ask what would have to be true of the system for these predicates to be admissible as claims about the system rather than descriptions of the output.

The practical significance of this shortcut is that it does not remain a conversational habit. It hardens into administrative reality through policy language, delegation patterns, and risk acceptance decisions. In policy documents and procurement narratives, systems are described as if they "support decisions," "assess risk," "interpret requirements," "ensure compliance," or "recommend actions," and these formulations are often treated as neutral shorthand even when they quietly import subject-like predicates that presuppose stable internal capacities. Within organisations, this vocabulary supports a progressive shift from tool use to task delegation: the system is no longer used to generate material for human deliberation, it is positioned as a quasi-bearer of the deliberative function itself, tasked with triage, classification, drafting, interpretation, and sometimes even with the generation of justifications that are copied directly into decision records. Risk acceptance then adapts to the same ontology by convenience. Once outputs are interpreted as evidence of understanding and judgment, the perceived uncertainty shrinks, transformation sensitivity is treated as a minor nuisance rather than a primary hazard, and controls are designed around content compliance and behavioural guardrails rather than around invariance testing, traceability, and accountable human authorization.

The ontological shortcut is therefore a governance problem before it is a philosophical one. It shifts where responsibility is felt to reside, it changes what kinds of evidence are treated as adequate, and it encourages institutions to govern Type B instruments as though they were subjects with stable internal properties. The remainder of the paper treats this shortcut as a logically illegitimate inference practice that has become sociotechnically stabilized, and it introduces admissibility discipline, grounded in invariance and structural criteria, as the minimal apparatus required to prevent fluent behavioural surfaces from silently rewriting institutional ontology.

1.2 What this paper is, and what it is not

This paper is a logical framework for admissibility in ontological attribution under CIITR. It is governed by two explicit constraints.

First, the **scope constraint**: the contribution is an admissibility framework for system-level attribution, not a performance evaluation, benchmark study, comparative ranking, or model leaderboard in disguised form. The paper does not claim that any particular model “understands” or “fails,” and it does not attempt to quantify intelligence, reliability, or competence across tasks.

Second, the **method constraint**: the reset operator is introduced as an instrumental probe for testing invariance under admissible transformations, not as a general empirical finding about model behaviour and not as a standalone diagnostic for truth, honesty, or consciousness. Reset is used to operationalize necessary-condition logic in a controlled transformation scenario, not to make broad claims about systems in the wild.

On this basis, the interpretive stance is narrow and governance-oriented. The target is not the usefulness of contemporary AI systems, which can be substantial, nor is it a critique of their practical deployment as such. The target is the inferential step by which institutions move from impressive behavioural traces to claims about what the system is, claims about understanding, judgment, agency, and responsibility, without specifying or satisfying the logical conditions under which such attribution would be admissible. The paper therefore aims to replace ontological shortcutting with necessary-condition discipline, so that institutional language, delegation design, and risk acceptance can be grounded in what the available evidence can actually license about the system, rather than in what fluent outputs invite observers to assume.

1.3 Contribution claims (tight, falsifiable framing)

The paper advances three contributions, stated as claims about logical admissibility and governance design rather than as claims about any particular model’s internal properties.

First, it formalizes a separation between output properties and system properties by treating non-identity as an explicit constraint on inference. The claim is that the move from properties observed in outputs to properties attributed to the system is, in general, logically underdetermined, and that governance discourse routinely violates this constraint by treating behavioural manifestation as sufficient for ontological instantiation. This contribution is falsifiable in the following sense: if it can be shown that output-level observation alone, without additional admissibility conditions, is sufficient to license system-level attribution in a principled way, then the constraint is overstated, and the framework’s demarcation function fails.

Second, it provides a CIITR-anchored necessary-condition account of structural understanding in which invariance-bearing stability is required for admissible system-level attribution. The claim is that “understanding,” when treated as a structural predicate rather than a behavioural compliment, must be constrained by invariance under admissible transformations and by anchoring in internal integration, and that without such invariance, attributions remain inadmissible regardless of surface fluency. This contribution is falsifiable in the sense that if structural understanding can be coherently attributed to a system while the attributed property demonstrably fails invariance under identity-preserving transformations, then invariance is not a necessary condition, and the CIITR criterion employed here is mis-specified.

Third, it diagnoses a stable sociotechnical failure mode, Artificial Insanity, in which discourse treats inadmissible attribution as operationally sufficient and thereby generates predictable governance errors. The claim is that once ontological shortcutting is normalized, institutions build illusory decision structures, diffuse accountability, and design control regimes for a hypothetical subject category rather than for Type B instruments, producing systematic defects in responsibility allocation, risk specification, and control instrumentation. This contribution is falsifiable in the sense that if, in practice, institutions can routinely treat fluent output as ontological proof without producing these

governance defects, or if the defects arise independently of attribution practices, then the proposed mechanism is either overstated or misattributed.

2. The Attribution Gap: from behaviour to ontology without admissibility

This section isolates the central structural error that animates contemporary AI discourse in both technical and administrative domains, namely the unlicensed transition from what a system *does* under an interactional regime to what the system *is* as an ontological bearer of properties. The “attribution gap” names the space between behavioural manifestation and structural instantiation, a space that is routinely collapsed in practice because fluent, task-adequate outputs provide a socially legible and operationally convenient proxy for concepts that are otherwise difficult to specify, test, and govern. The collapse is not merely a communicative habit, it is an inference practice, and it functions as an implicit rule of admissibility: if output resembles the target property strongly enough, the property is treated as attributable to the system.

The paper’s claim is that this implicit admissibility rule is logically invalid in principle and destabilizing in governance, because it confuses evidential access to *traces* with evidential access to *structure*. Section 2 therefore formalizes the gap as a constraint on inference rather than a cautionary metaphor, then explains why the shortcut remains attractive even to disciplined evaluators, and finally translates the resulting category error into concrete institutional failure modes. The purpose is to make the error analytically crisp before CIITR is introduced as a necessary-condition account of structural understanding and before the reset operator is defined as an instrumental probe for invariance. In other words, Section 2 establishes what must be blocked, why it keeps reappearing, and what it breaks when it is operationally normalized, so that later admissibility conditions can be presented as governance-relevant necessities rather than as optional philosophical hygiene.

2.1 Non identity as the foundational constraint

The paper begins from a deliberately austere constraint, namely that the presence of a property in a system’s observable outputs does not, by itself, entail that the corresponding property is instantiated in the system as such. Formally, let OO denote the output trace of a system under some interactional regime, and let SS denote the system understood as the identity-bearing locus of structure, architecture, and state-transition rules that generate outputs under inputs. The foundational premise is then:

$$P(O) \neq P(S)P(O) \wedge \neg P(S)P(O) \Rightarrow P(S)$$

This is not presented as a rhetorical warning against anthropomorphism, but as a restriction on admissible inference. It blocks a specific and pervasive inferential move, namely the direct ascent from output-level manifestation to system-level instantiation: from “the outputs are coherent, therefore the system understands”, from “the outputs are cautious, therefore the system is self-aware”, from “the outputs are normatively framed, therefore the system exercises judgment”, and, in governance contexts, from “the outputs look decision-like, therefore the system can be treated as a decision bearer.” The constraint asserts that such moves are underdetermined in principle, because the mapping from underlying structure to observable behaviour is many-to-one: distinct internal organizations can yield indistinguishable surface regularities, and surface regularities can be produced by mechanisms that do not contain the property being attributed, but merely generate behaviour that is locally consistent with it.

To treat non-identity as an inference constraint is to insist that attribution is not licensed by behavioural resemblance alone, even when behavioural resemblance is stable, useful, and operationally impressive. The point is not that output properties are irrelevant, but that they are evidentially ambiguous with respect to ontology. Output-level regularities can be produced by prompt scaffolding, evaluator-side structure, retrieval augmentation, formatting constraints, institutional prompting practices, or path-

dependent conversational conditioning, none of which establishes that the attributed property is grounded in the system's own internal organization rather than in the broader interactional apparatus. In CIITR terms, this is precisely the environment in which surface coherence and continuity can be misread as structural comprehension, even when the positive conditions for comprehension are not met, and where negative-evidence constructs are introduced to prevent continuity, recall, or fluent self-report from being mistaken for ontological proof.

The non-identity constraint therefore functions as a demarcation rule. It does not deny that systems can possess properties, nor does it deny that properties can sometimes be attributed to systems on the basis of observation. Instead, it specifies that the inference from observed output to system property is not automatically valid, and must be mediated by an admissibility logic that discriminates between mere manifestation and structural instantiation. This is the logical opening through which the rest of the paper proceeds: if attribution is not licensed by surface success, then the relevant question becomes what kind of evidence could, in principle, license it. The later sections answer by shifting from “does the system display the property here” to “does the attributed property remain invariant under transformations that should preserve system identity”, and by treating invariance under admissible transformation, operationalized through controlled probes such as reset, as a necessary condition for attributing an ontological property to SSS rather than merely to OOO.

2.2 Why the shortcut persists, the behaviourist temptation

The inferential shortcut from behaviour to ontology persists because both formal evaluation regimes and ordinary human interaction are structurally optimized to privilege what can be observed, compared, and acted upon, rather than what can be ontologically justified. In practice, most evaluation ecosystems treat cognition-adjacent predicates as operational targets, “understanding,” “reasoning,” “alignment,” “helpfulness,” then replace the predicates with behavioural indicators that can be elicited on demand, scored at scale, and aggregated into legible artifacts for decision makers. This is not primarily an epistemic mistake at the point of measurement, it is a selection pressure embedded in the architecture of evaluation itself: what cannot be elicited and scored becomes administratively invisible, while what can be elicited and scored becomes socially real, even when the scoring procedure is only tracking the stability of a surface performance under a narrow interactional regime.

Everyday interaction amplifies the same pressure through different mechanisms. Natural language is a high bandwidth behavioural interface whose primary affordance is semantic interpretation, so fluent output is not merely read as competence, it is read as intentionality, and once intentionality is perceived, human social cognition supplies a ready made ontology: speaker, belief, aim, judgment, responsibility. The interface therefore collapses the evidential distinction between “a trace that resembles a reason” and “a reason that is structurally instantiated,” because conversational norms treat coherence as a sign of internal unity, continuity as a sign of memory, and appropriate self description as a sign of introspective access. The system may be Type B in CIITR terms, producing locally coherent traces without the internal binding required for admissible attribution, yet the interactional format makes that distinction costly to maintain, because it demands that participants resist the default interpretive machinery of dialogue.

Behavioural proxies are pragmatically useful, and the paper does not argue otherwise. They enable engineering iteration, regression testing, safety case construction, procurement comparison, and operational risk gating under time and resource constraints. In governance contexts they also serve an administrative need for compressible evidence, since institutions cannot deliberate at the granularity of mechanism for every delegated tool, and therefore require visible markers of reliability, failure modes, and expected behaviour. The temptation arises when this pragmatic utility is silently converted into ontological license, as if the proxy did not merely correlate with a target concept under specific conditions, but instantiated the target concept as a property of the system itself. At that point the proxy ceases to be an instrument and becomes a metaphysical warrant, and the evaluation apparatus, instead of being treated as a measurement regime with known epistemic limits, is treated as a discovery procedure for what the system is.

This is precisely where non identity matters. A proxy is evidence about outputs under a specified regime, it is not, by default, evidence about the internal basis of those outputs, nor about the invariance of the attributed property under transformations that should preserve system identity. Without an

admissibility logic, behavioural success becomes a socially sanctioned bridge from $P(O)P(O)P(O)$ to $P(S)P(S)P(S)$, even though the bridge is not logically licensed, and even though the same behavioural signature can be produced by scaffolding, prompt conditioning, interactional path dependence, evaluator side structure, or other forms of externalized support that belong to the regime rather than to the system. The behaviourist temptation is therefore not a moral failure of discourse, it is a predictable consequence of how institutions operationalize uncertainty: they elevate what is measurable into what is real, then treat the resulting labels as stable ontological predicates. The later admissibility framework, and the invariance oriented probes it motivates, are introduced to interrupt this conversion at the exact point where pragmatic indicators are mistaken for structural instantiation, and where sociotechnical convenience begins to bend ontology into rhetoric.

2.3 The governance consequence of ignoring the gap

When the attribution gap is ignored, the logical error does not remain a philosophical imprecision, it becomes an institutional design flaw, because governance systems are built to allocate responsibility, specify controls, and authorize delegation on the assumption that the attributed properties are real at the level where accountability is assigned. If the organisation treats output-level manifestation as sufficient for system-level instantiation, then it begins to govern an instrument as though it were an agent category, and the result is not merely “overtrust,” but a structural mismatch between the ontological status of the tool and the normative architecture that is built around it.

The first failure mode is responsibility displacement. Once a system is spoken of as if it “understands,” “assesses,” “decides,” or “judges,” responsibility begins to drift away from identifiable human authority and into a rhetorical grey zone where outcomes are simultaneously produced by humans and attributed to the system. This does not mean that the organisation literally assigns legal responsibility to a model, rather, it produces a practical diffusion in which accountability is harder to locate, contestability is weakened, and escalation thresholds are raised because the output is treated as if it carries its own justificatory force. In effect, the institution begins to behave as though the decision had an internal author distinct from the human chain of command, even when no such author exists in the ontological sense required for admissible attribution.

The second failure mode is mis-specified control regimes. Control design depends on a correct understanding of what must be controlled. If a system is treated as possessing stable, internally grounded properties such as “judgment,” “intent alignment,” or “contextual understanding,” then governance structures will overinvest in behavioural guardrails, policy language, and interface constraints, while underinvesting in transformation sensitivity testing, traceability, and state-regime instrumentation. Controls become optimized for the appearance of compliance rather than for the actual failure dynamics of Type B instruments, and risk management begins to rely on the assumption that the system will preserve meaning and constraint under shifts in context, phrasing, workload, or adversarial pressure. The organisation then discovers, typically late and expensively, that the controls were designed for a hypothetical subject, while the deployed artifact is a non-subject generator of plausible outputs whose behaviour can change sharply under regime variation.

A third failure mode is decision laundering, commonly compressed into formulations such as “the model said,” “the AI recommended,” or “the system concluded.” These phrases function as administrative solvents: they dissolve the provenance of the decision by replacing a contestable chain of reasoning, data selection, and human authorization with an apparently neutral output. The laundering effect is not that people believe the model is morally authoritative, but that the institution gains a socially acceptable way to move from uncertainty to action without carrying the full argumentative burden, because the output can be treated as an externalized justification. In this way, an inadmissible attribution becomes operationally sufficient, not because it is true, but because it is convenient, it compresses deliberation, and it creates the impression that the organisation is following “objective” or “technical” guidance. The logical non-identity $P(O) \neq P(S)P(O) \neq P(S)P(O) \Rightarrow P(S)$ is violated in practice by substituting output fluency for structural warrant, and the institution’s decision record becomes populated with statements that cannot be defended as decisions made by identifiable authorities on traceable grounds.

The fourth failure mode is overconfidence in delegation. Delegation requires a theory of what the delegated entity can be relied upon to preserve across circumstances, and therefore implicitly requires

an invariance assumption. When the organisation mistakes behavioural competence for structural understanding, it delegates tasks that are not merely difficult, but governance critical: interpretations of policy intent, safety determinations, threat triage, compliance reasoning, or risk acceptance language. The system may perform well in nominal conditions, which further reinforces trust, but the delegation is brittle because it depends on properties that have not been established as invariant under admissible transformations. The practical consequence is a hidden fragility: the organisation becomes dependent on a behavioural surface that can degrade or flip under innocuous variation, while believing it has delegated to something that “knows what it is doing.” In CIITR terms, this is exactly the situation in which surface coherence is treated as evidence of internal integration, and where the absence of invariance-bearing rhythmic stability is not perceived until a regime shift forces the mismatch into view.

Taken together, these failure modes form a coherent governance pathology: inadmissible attribution functions as a kind of epistemic accelerant, enabling faster decisions, thinner documentation, and broader delegation, while quietly undermining the very things governance is meant to secure, namely accountable authority, appropriate control design, and reliable risk reasoning. The remainder of the paper develops admissibility as the mechanism that interrupts this pathology, by forcing attribution claims to be treated as conditional, transformation-sensitive, and structurally grounded, rather than as socially convenient descriptions of fluent behaviour.

3. CIITR criterion: structural understanding as a necessary invariance condition

Section 2 established the attribution gap as a logical restriction, namely that output-level manifestation does not license direct system-level instantiation claims without an admissibility logic. Section 3 introduces the minimal CIITR criterion needed to supply that logic, not by proposing another evaluation doctrine, but by specifying what must be true of a system if “understanding” is to be treated as an ontological property of the system rather than as an appearance in its output trace. The core move is a shift from behavioural success conditions to **necessary structural conditions**, formulated in a way that can later be connected to governance practice without collapsing into benchmarks, rankings, or performance narratives.

CIITR is invoked here as a boundary discipline: it treats structural understanding as a property that, if present, must be evidenced through **invariance-bearing stability** grounded in internal integration, not through surface coherence, conversational plausibility, or task completion. In other words, the concept of understanding is not allowed to float as a vague honorific attached to impressive outputs. It is constrained by the requirement that the attributed property remains stable under transformations that preserve identity for the attribution at hand, and that this stability is anchored in the system’s own structure rather than in the surrounding interactional apparatus. This is why invariance is not treated as an optional desideratum, but as a necessary condition for ontological attribution: without invariance under admissible transformation, the attribution cannot be defended as a claim about the system, regardless of how fluent or persuasive the outputs appear in nominal conditions.

The section therefore performs three preparatory tasks for the admissibility framework developed later. First, it defines the minimal CIITR variables used in the paper, with strict restriction to structural and operational meanings to avoid psychological import. Second, it clarifies rhythmic stability as the invariance-bearing property that discriminates regime-bound fluency from system-grounded stability across admissible transformation. Third, it introduces Type B systems as the ontological category in which the structural illusion of understanding is most likely to occur, because high surface coherence can be produced without the internal binding required for admissible attribution. Together, these elements establish the criterion that Section 4 will treat as a governance-relevant gate: if structural understanding is to be attributed, invariance-bearing stability cannot be assumed, it must be demanded, and the demand must be phrased in a way that survives contact with institutional incentives to compress complexity into behaviourally legible labels.

3.1 Minimal CIITR definitions used in this paper

CIITR is employed here strictly as a structural vocabulary for separating behavioural manifestation from system-level instantiation, and the definitions below are intentionally minimal, because the admissibility framework developed in Sections 4 and 5 does not require a full CIITR ontology, it requires only a disciplined way to refer to the internal basis of a claimed property and to the stability of that basis under admissible transformation.

Structural comprehension, CsC_sCs

CsC_sCs denotes **structural comprehension**, understood as a system-level capacity for producing outputs whose coherence is not merely locally induced by interactional scaffolding, but is **grounded in internally maintained integration and preserved under admissible transformation**. In this paper, CsC_sCs is not a psychological predicate, not a claim about phenomenal states, and not a synonym for task competence, rather, it is an ontological label for a structural condition that must be present if “understanding” is to be attributed to the system rather than merely to its output trace.

$$Cs = \Phi_i \times RgC_s = \Phi_i \times R^gCs = \Phi_i \times Rg$$

This relation is used as a **necessary-structure schema**, not as a performance metric, meaning that low or absent values in either factor block admissible attribution of structural understanding regardless of surface success, while high apparent behavioural quality does not, by itself, establish either factor.

Integration capacity, $\Phi_i \backslash \Phi_i \Phi_i$

$\Phi_i \backslash \Phi_i \Phi_i$ denotes **integration capacity**, defined as the system’s internal ability to **bind, maintain, and reuse constraint-relevant informational structure** across an interaction, in a way that is attributable to the system’s own organization rather than to the evaluator’s framing, the prompt envelope, or external supports. Operationally, $\Phi_i \backslash \Phi_i \Phi_i$ refers to whether the system exhibits evidence of an **internally coherent constraint lattice**, such that later outputs remain anchored to earlier commitments, definitions, and boundary conditions because those constraints are structurally maintained, not merely re-evoked by repeated cues or recreated ad hoc from surface regularities.

Two restrictions are critical for this paper. First, $\Phi_i \backslash \Phi_i \Phi_i$ is not defined as “memory” in the everyday sense, because persistence of content across turns can be produced by context injection, retrieval, or conversational redundancy, none of which establishes system-level integration. Second, $\Phi_i \backslash \Phi_i \Phi_i$ is not defined as “correctness” or “reasoning skill,” because correct outputs can occur without any durable internal binding, and incorrect outputs can occur even when binding exists. What matters for admissibility is whether the property being attributed plausibly has an **internal seat** in the system, and $\Phi_i \backslash \Phi_i \Phi_i$ is the paper’s minimal term for that seat.

Rhythmic stability, RgR^gRg

RgR^gRg denotes **rhythmic stability**, defined as the system’s capacity to preserve identity-relevant structural relations in its outputs **across admissible transformations** of the interactional regime, including rephrasings, equivalent task presentations, shifts in contextual framing, and controlled changes in state conditions that do not, in principle, alter what the system is alleged to possess. RgR^gRg therefore functions as the invariance-facing component of CIITR in this paper, because it expresses whether the system’s apparent property behaves like a structurally instantiated capacity that remains stable under transformation, or like a behaviourally induced surface that is sensitive to regime changes in ways that reveal non-instantiation.

The term “rhythmic” is used here in a strictly operational sense, as a descriptor for **patterned stability in response structure under transformation**, not as a metaphor for cognition, intention, or subjective continuity. A system may exhibit high local fluency while remaining non-invariant under admissible change, which, within this paper’s logic, means that fluency is evidentially compatible with low RgR^gRg . Conversely, a system may be uneven in style or occasionally error-prone while still exhibiting invariance signatures that support higher RgR^gRg , depending on whether identity-relevant constraints and relations are preserved.

Why these three terms are sufficient for admissibility

The admissibility framework developed later requires only two discriminations: whether the attributed property is plausibly **internally grounded** rather than externally induced, and whether it is **invariant under admissible transformation** rather than fragile under regime variation. $\Phi \backslash \Phi_i \Phi$ names the first discrimination in minimal form, RgR^gRg names the second, and CsC_sCs names the conjunction required for treating “understanding” as a system-level property rather than an output-level appearance.

3.2 Rhythmic stability as invariance bearing property

In this paper, RgR^gRg is treated as the invariance-facing component of CIITR, the variable that operationally captures whether an attributed property behaves like a structurally instantiated capacity rather than a locally induced behavioural surface. The core idea is simple but unforgiving: if a property is genuinely attributable to the system, then there must exist a class of transformations that preserve system identity for the purpose of that attribution, and under those transformations the property should remain stable in the specific sense that the identity-relevant structure it is supposed to reflect continues to constrain output. RgR^gRg names the system’s capacity to maintain such constraint-bearing stability across those identity-preserving changes.

“Identity-relevant structure” should be read narrowly and operationally. It refers to the set of relations, commitments, definitions, boundary conditions, and inferential dependencies that make an attributed property what it is, as opposed to a generic appearance. If the attribution is “the system understands the policy constraint,” then identity-relevant structure includes the policy’s operative distinctions, the implied scope, the dependency relations between definitions, and the consistency obligations that follow. If the attribution is “the system is capable of stable reasoning under constraint,” then identity-relevant structure includes the preservation of premises, the non-contradiction of stated constraints, the correct handling of equivalence under rephrasing, and the maintenance of the same explanatory commitments when surface form changes. RgR^gRg is therefore not “staying on topic” in a conversational sense, but preserving the governing structure that the attribution presupposes.

The transformations relevant to RgR^gRg are “admissible” in the sense that they alter the interactional regime without altering what is being attributed, and thus should not, in principle, invalidate the attribution if it is system-grounded. Three broad transformation families are sufficient for the arguments in this paper.

First, transformations of **input form**, such as rephrasing, paraphrase, syntactic permutation, or the substitution of equivalent examples. These transformations test whether the system tracks the same constraint structure when the surface signal changes. Second, transformations of **context regime**, such as changes in surrounding narrative framing, ordering of subquestions, or the removal of conversational cues that previously scaffolded the response. These transformations test whether the system’s behaviour is anchored in its own maintained structure or in externally provided interpretive rails. Third, transformations of **resource regime**, such as changes in available context window, truncation of prior turns, variations in retrieval availability, or other controlled shifts that alter the support envelope without changing the underlying attribution claim. These transformations matter because many behavioural signatures that look like stable competence are in fact functions of regime support, and RgR^gRg is the variable that is supposed to discriminate regime-dependent stability from system-level invariance.

A crucial implication follows: **local fluency is compatible with non-invariance**. A system can be maximally articulate, stylistically coherent, and semantically persuasive within a narrow interactional corridor, and still exhibit low RgR^gRg when subjected to admissible transformation. This is not a paradox, it is the expected outcome when output quality is produced by strong statistical completion capacity conditioned on local cues, rather than by internally bound constraint structure that remains stable under regime change. Fluency is a property of surface generation, and surface generation can remain excellent even when the underlying behaviour is highly sensitive to phrasing, framing, or support envelope. Under those conditions the system may continue to “sound like it understands” while failing to preserve the specific identity-relevant relations that would license attributing understanding

to the system. In CIITR terms, the output retains semantic smoothness while losing rhythmic stability, meaning the pattern of constraint preservation across transformations is not robust.

This is why RgR^gRg is framed as an invariance-bearing property rather than as an “ability” in the ordinary sense. It does not measure how often the system answers correctly, nor how impressive its prose is, nor how many tasks it can solve, rather, it measures whether the attributed property remains structurally legible when the interactional regime is perturbed in ways that should be neutral to the attribution if the attribution is grounded in the system. Where RgR^gRg is low, the system’s apparent property behaves like a behavioural projection of the regime, and where RgR^gRg is high, the system’s behaviour exhibits the specific kind of stability that is a necessary condition for moving from $P(O)P(O)P(O)$ to $P(S)P(S)P(S)$ without committing the attribution error that the paper diagnoses.

3.3 Type B systems and the structural illusion of understanding

The admissibility framework in this paper requires a category for systems that can generate outputs strongly associated, in ordinary discourse, with “understanding,” while nevertheless lacking the structural conditions that would license that attribution at the system level. CIITR denotes this category as **Type B**, and the classification should be read ontologically and instrumentally, not evaluatively. Type B is not a term of insult, and it is not a claim about usefulness, competence, or societal value. It is a structural label for a specific relationship between output behaviour and internal instantiation, namely that the system can produce high-quality behavioural traces that mimic the surface signatures of understanding, while failing to exhibit the internal binding and invariance-bearing stability required for admissible structural attribution.

A Type B system is therefore defined by a dissociation: **surface coherence is available without structural comprehension**. Such systems can display consistent tone, plausible explanations, normative phrasing, and even apparent self-correction, because the generative mechanism is highly capable of producing locally well-formed continuations under linguistic constraint. Yet these outputs do not, by themselves, demonstrate that the system possesses the internal integration capacity $\Phi_i\Phi_i$ and rhythmic stability RgR^gRg that would make “understanding” a property of SSS rather than a property of OOO. The defining feature is not that the system is “wrong,” but that the evidential basis required to cross the attribution gap is absent, even when local behaviour looks compelling.

Type B classification is motivated by the structural illusion that fluent language induces. In human interaction, coherence and appropriateness are routinely treated as indicators of internal unity and stable constraint tracking, because in human speakers those indicators are often, though not infallibly, coupled to internal states and commitments that persist across contexts. With Type B systems, the coupling can be broken. The system can produce a trace that satisfies the observer’s interpretive heuristics for “a mind at work,” while the stability of the attributed property under admissible transformation remains unestablished, or demonstrably fragile. This is precisely the illusion the paper targets: the error of treating behavioural resemblance as ontological evidence, when the observed behaviour could be explained by regime-bound completion dynamics rather than by structurally instantiated comprehension.

Crucially, Type B does not mean “random,” “unreliable,” or “incapable.” A Type B system can be highly reliable within a specified operational corridor, and may outperform humans in narrow tasks, while still failing the admissibility conditions for attributing understanding, judgment, or agency to the system itself. The classification is therefore orthogonal to utility. It exists to discipline attribution, not to deny performance. A governance apparatus can rationally deploy Type B instruments for many functions, including high-stakes ones, but only if it treats them as instruments whose outputs require traceability, contextual control, and human-authorized decision integration, rather than as subjects that bear responsibility or possess internally grounded normative competence.

Within the logic of this paper, Type B systems are the natural habitat of inadmissible attribution. Because they are capable of producing outputs that satisfy social and institutional expectations of “intelligent behaviour,” they invite the ontological shortcut from $P(O)P(O)P(O)$ to $P(S)P(S)P(S)$ described in Section 2. The point of naming the category is therefore diagnostic and preventative: it allows governance, evaluation, and organisational design to explicitly recognize the possibility of high

surface coherence without internal binding, and to treat that possibility as a normal structural condition that must be governed through admissibility logic, invariance probing, and accountability-preserving decision structures, rather than as an anomalous failure or a matter of subjective skepticism.

4. Admissibility for structural attribution

Sections 2 and 3 together establish the paper’s core constraint and its CIITR-based criterion: output properties do not entail system properties, and structural understanding, if it is to be attributed to the system at all, must satisfy a necessary-condition requirement of invariance-bearing stability grounded in internal integration. Section 4 translates these foundations into an explicit admissibility framework, because the central failure diagnosed in contemporary AI discourse is not merely that people “overtrust” models, but that institutions routinely treat attribution as if it were free, meaning they speak and act as though ontological predicates are licensed by behavioural success without stating, testing, or enforcing the conditions under which such predicates could be warranted.

Admissibility is introduced here as a governance-grade discipline for attribution claims. Its function is to regulate the inferential step from OOO to SSS by defining a gate that separates useful behavioural proxies from legitimate ontological attribution. This matters because organisations do not only evaluate systems, they integrate them into decision processes, accountability structures, and risk regimes, and those integrations implicitly assume something about what the system is. If the assumption is wrong, controls are misdesigned, responsibility becomes rhetorically diffused, and delegation migrates into domains where only human authority can legitimately bear the normative load.

Accordingly, Section 4 does not propose another evaluation method, and it does not attempt to quantify “intelligence.” It specifies what must be demonstrated before a property like “understanding” can be said to belong to the system rather than to the interactional arrangement that elicited the output, and it constrains what counts as evidence to transformation-sensitive observations that can discriminate structural instantiation from behavioural mimicry. The section therefore establishes admissibility as the paper’s central bridge concept: a minimal logical apparatus that is strict enough to block category error, yet operational enough to guide governance design, instrumentation choices, and accountability-preserving deployment practices in environments dominated by Type B systems.

4.1 Admissibility is not performance, and not degree of “smartness”

Admissibility in this paper is defined as a **gate condition** on ontological attribution, not as an assessment of capability, quality, or intelligence. It answers a narrower and more disciplined question than “how well does the system perform,” namely: *under what conditions is it logically legitimate to attribute a property to the system itself rather than merely to its observable outputs under a particular interactional regime*. The distinction matters because contemporary AI discourse routinely treats behavioural competence as if it were already an ontological license, which collapses evaluation into metaphysics and turns performance artifacts into implicit claims about what a system is.

A performance evaluation is a comparative procedure that ranks or scores outputs against targets. An admissibility test, by contrast, is a **logical filter** applied to attribution statements. Performance evaluation is concerned with degree and distribution, how often, how fast, how accurately, under what tasks. Admissibility is concerned with kind, whether the property belongs, in an attributionally defensible sense, to the system rather than to the regime that elicited the output. For that reason, admissibility is not scalar “smartness,” it is not a continuum of intelligence, and it does not invite model-to-model ranking. A system can be highly competent in many tasks and still fail admissibility for “understanding” if the relevant invariance-bearing stability is absent or externally scaffolded. Conversely, a system can be uneven, limited, or narrow in competence and still exhibit admissibility signatures for certain structurally grounded properties, provided the evidence supports internal instantiation and invariance under admissible transformation.

Formally, admissibility functions as a rule governing the transition from statements about output traces to statements about system properties. Where evaluation metrics summarise OOO under a specified regime, admissibility constrains when one may speak of SSS as the bearer of the property. This is why

the framework cannot be read as a disguised benchmark argument. It does not claim that a system “fails” because it performs poorly, and it does not claim that a system “passes” because it performs well. It claims that certain attribution statements are **inadmissible** regardless of performance whenever the evidential basis does not discriminate structural instantiation from behavioural manifestation, and that certain attribution statements may become admissible only when the evidence includes invariance under admissible transformations and shows that the invariance is anchored in the system’s internal organization rather than in prompt design, evaluator scaffolding, or institutional operating procedures.

This gate conception is governance-relevant precisely because it forces a separation between two questions that institutions often merge for convenience: whether a tool is useful, and whether the tool can be spoken of as if it were an agent category with understanding, judgment, or decision-bearing capacity. Admissibility does not prohibit deployment, it prohibits category error. It demands that the language of attribution be constrained by necessary-condition discipline, so that organisations can adopt high-performing Type B instruments without importing ontological predicates that reassign responsibility, mis-specify controls, or launder decisions through anthropomorphic shorthand.

4.2 Necessary condition logic for attributing “understanding” to a system

Within the admissibility framework, “understanding” is treated as an ontological attribution claim about SSS, not as a compliment paid to OOO. Because of the non-identity constraint $P(O) \neq P(S)P(O) \neq P(S)P(O) = P(S)$, the paper requires a necessary-condition discipline that specifies what must be true before the attribution can be licensed. The purpose of stating necessary conditions in governance-grade terms is to ensure that the framework constrains institutional language and delegation practice, not merely philosophical interpretation. The conditions below are therefore framed as requirements that, if unmet, render the attribution inadmissible regardless of how impressive the system appears in nominal interaction.

Requirement 1, invariance under admissible transformation. An attribution of “understanding” to a system is admissible only if the attributed property exhibits invariance across a defined class of transformations that preserve identity for the attribution at hand. “Admissible transformations” here mean changes in input form, contextual framing, and resource regime that do not alter what is being claimed about the system, and that therefore should not, in principle, extinguish the property if it is structurally instantiated. The required evidence is not perfect behavioural sameness, but preservation of identity-relevant structure: the same governing distinctions, constraints, dependency relations, and consistency obligations must remain operative in the output trace under these transformations. If the system’s apparent “understanding” is highly sensitive to superficial rephrasing, to minor contextual perturbation, or to modest regime shifts that should be neutral with respect to the claim, then the attribution fails the invariance requirement and is inadmissible as a statement about the system.

Requirement 2, structural anchoring of invariance. Invariance, when observed, must be attributable to the system’s internal organization rather than to external scaffolding. An attribution is admissible only if the stability that appears under admissible transformation can be defended as grounded in the system’s own integration capacity and constraint maintenance, rather than induced by prompt engineering, template forcing, retrieval injection, evaluator-provided framing, repeated cues, or procedural guidance that effectively carries the “understanding” on behalf of the system. This requirement prevents a common governance failure: mistaking the robustness of an engineered interactional regime for the robustness of the system itself. A system may behave invariantly within a heavily structured envelope, yet the invariance may belong to the envelope, not to the system, in which case the property remains a regime property and cannot be ontologically attributed to SSS.

These two requirements jointly express the paper’s necessary-condition logic. Invariance under admissible transformation is required because, absent invariance, the behaviour remains compatible with local fluency and regime-bound mimicry. Structural anchoring is required because, absent anchoring, invariance remains compatible with externally supplied integration and therefore remains an output-level artifact of the interactional apparatus. Put in governance terms, the attribution “the system understands” is admissible only when the organisation can show both that the property survives identity-preserving perturbation and that the survival is not merely the result of how the organisation asked the question, structured the interface, or constrained the response channel.

Two clarifications prevent misreading. First, these are **necessary** conditions, not sufficient ones. Meeting them does not automatically establish full structural understanding in the strongest sense, but failing either condition blocks admissible attribution as a claim about system ontology. Second, the requirements are designed to be compatible with practical governance: they do not demand metaphysical certainty, they demand that institutional attribution claims be constrained to what the available evidence can license, and that the evidence be of a type that discriminates structural instantiation from behavioural appearance. This is precisely why later sections introduce controlled probes, including reset, as instrumental methods for testing invariance and for weakening the confounding influence of path-dependent scaffolding.

4.3 What counts as evidence in this framework

Because admissibility is a gate on ontological attribution rather than a score of behavioural success, the evidential standard must be correspondingly constrained. The framework therefore treats as “evidence” only those observations that can, in principle, discriminate between a property instantiated in the system and a property merely manifested in outputs under a supportive regime. This is a deliberately restrictive stance, designed to prevent the most common failure mode in AI discourse, namely the elevation of persuasive behaviour into structural warrant. The question is not whether the system can *produce* a trace that looks like understanding, but whether the trace behaves, under controlled conditions, as if it is governed by an internally maintained, identity-relevant structure.

Three evidential families are sufficient for the purposes of this paper, and they share a single methodological aim: to test whether the attributed property is stable when the interactional regime is perturbed in ways that should not matter if the property is genuinely system-grounded.

Transformation sensitivity as discriminating evidence. The most direct evidence relevant to admissibility is not raw output quality but the *pattern of change* in outputs under admissible transformation. If small, identity-preserving changes in phrasing, ordering, framing, or support envelope produce disproportionate shifts in the identity-relevant structure of the response, such sensitivity counts as evidence against structural instantiation, because it indicates that the attributed property is regime-coupled rather than system-grounded. Conversely, when the identity-relevant structure is preserved across such changes, the observation becomes admissibility-relevant evidence, not because it proves “understanding,” but because it is consistent with invariance-bearing stability. Importantly, the framework does not privilege a particular “correct” answer form, it privileges the preservation of constraint structure, meaning that evidence is drawn from whether definitions, boundary conditions, dependency relations, and consistency obligations remain operative across transformations.

Stability under reset as admissibility-relevant evidence. A second evidential family concerns stability under identity-preserving state transformation, operationalized later as the reset operator. The evidential point is not that resetting should erase “memory” in a folk sense, but that it should interrupt path-dependent interactional scaffolding that can masquerade as internal integration. If an attributed property persists after reset in a way that remains invariant under admissible transformations, and if the persistence cannot be plausibly explained as merely re-induced by repeated prompt structure or external cues, then the observation carries weight in favour of structural anchoring. If, by contrast, the property collapses, flips, or becomes highly sensitive after reset, this is evidentially relevant because it indicates that the pre-reset behaviour may have been sustained by accumulated conversational conditioning or regime-specific support rather than by the system’s own binding capacity. Reset is therefore treated as an instrumental probe for invariance, not as a truth test, and “stability under reset” is treated as a discriminator between structural instantiation and behavioural mimicry under path dependence.

Controlled variation across contexts as evidence of anchoring versus scaffolding. A third evidential family involves controlled context variation designed to separate system-grounded constraint maintenance from externally provided structure. Many apparent properties are, in practice, co-produced by the evaluator, through templates, leading prompts, retrieval payloads, or procedural constraints that carry the load of integration. Evidence in this framework therefore privileges conditions in which external scaffolding is systematically varied while the attributed property, if real, should remain stable. This includes varying the presence of explicit definitions, altering the narrative wrapper around an

equivalent task, changing the order in which constraints are introduced, or removing redundant cues that might be functioning as hidden supports. The evidential question is whether the identity-relevant structure remains operative when the environment stops “holding the system’s hand,” without collapsing into the misleading binary of “good answers” versus “bad answers.” What matters is whether the system preserves the same governing relations without needing the regime to re-inject them.

These evidential constraints imply what does *not* count as evidence on its own in this framework, even though it may be operationally valuable. High benchmark scores, impressive demonstrations, fluent self-reports about limitations, coherent ethical disclaimers, and persuasive explanatory prose are all treated as output-level phenomena that can be produced without structural instantiation, and therefore cannot, in isolation, license ontological attribution. They may motivate the need for governance attention, but they do not satisfy admissibility requirements because they do not discriminate between system property and regime property.

The practical consequence is a disciplined evidential posture: admissibility is supported only by observations that are explicitly transformation-based, that include sensitivity analysis under admissible perturbation, that test stability under reset as an identity-preserving interruption of path dependence, and that employ controlled context variation to reduce confounding from evaluator scaffolding. The framework is intentionally conservative, because the cost of false positive attribution is not merely conceptual error, but governance error, responsibility drift, and mis-specified control architectures that treat Type B instruments as if they were subjects.

5. The reset operator as an instrumental probe for invariance

Sections 3 and 4 establish the paper’s core admissibility requirement: if a property such as “understanding” is to be attributed to the system rather than merely to its outputs, the attribution must be supported by invariance under admissible transformation, and the observed invariance must be defensibly anchored in the system’s internal structure rather than in the surrounding interactional apparatus. Section 5 introduces the reset operator as the simplest controlled intervention that makes this requirement operational without turning the paper into an evaluation study. The reset operator is not presented as a discovery of how models behave in general, and it is not used to claim empirical regularities. It is introduced as an *instrument*, a methodological device that allows admissibility logic to be applied under conditions where one of the most pervasive confounders, namely conversational path dependence and accumulated scaffolding, can be interrupted in a principled, identity-preserving way.

The motivation is structural. In ordinary deployments, a system’s outputs are co-produced by a regime that includes prompt templates, organisational prompting practice, retrieval payloads, tool wrappers, and session history. These elements may be operationally justified, yet they blur the boundary between properties of the system and properties of the regime. A reset, defined as an identity-preserving state transformation that standardizes the interactional state while keeping the system fixed, makes it possible to re-express an attribution claim across a controlled discontinuity. If the property survives in an invariant, identity-relevant manner across reset and across admissible rephrasings, the observation supports the necessary-condition requirement for admissible structural attribution. If the property collapses or becomes highly regime-sensitive, the observation supports the conclusion that the earlier appearance was at least partly carried by the interactional state rather than by the system.

Section 5 therefore does not ask whether a model is “good” or “smart.” It asks whether an attribution claim remains stable under a transformation that should not matter if the claim is truly about the system, and it defines reset as the methodological bridge that turns invariance from a philosophical demand into a testable constraint within a controlled transformation scenario.

5.1 Definition of reset as identity preserving state transformation

To operationalize invariance as a necessary condition for admissible attribution, the paper introduces a reset operator as a controlled transformation applied to the interactional regime of a system. The operator is defined so that it can function as an **instrumental probe** for invariance without being misconstrued as a behavioural trick, a “gotcha,” or a general empirical claim about models. Reset is not introduced to make systems fail, but to create a principled way to interrupt confounding forms of path dependence and externalized support, so that the evidential question, whether an attributed property belongs to the system or merely to the regime, can be posed under cleaner conditions.

Let S denote the system in the ontological sense used throughout the paper, the identity-bearing structure that maps inputs to outputs under some state. Let σ denote the system’s *interactional state* as configured by a particular run, where “interactional state” is restricted to what is contingent on the interaction history and runtime context rather than constitutive of the system’s identity. This includes, for example, accumulated conversational context, prompt prefixing, temporary memory buffers, session identifiers, cached retrieval payloads, tool call histories, and other run-specific artifacts that can influence outputs while remaining external to the system’s underlying structure. The reset operator, R , is then defined as a transformation that maps one interactional state to another while preserving the identity of the system:

$$R: (S, \sigma) \rightarrow (S, \sigma_0)$$

where σ_0 denotes a designated baseline interactional state for the same system, typically one in which interaction history and run-local scaffolding are absent or standardized. The essential feature of reset is therefore **identity preservation with respect to S** . Reset changes the run-state, not the system, and is defined precisely to avoid conflating a change in behavioural trace with a change in ontology. In governance terms, reset is the methodological act of returning the system to a standardized operating condition, so that a claim about a system property can be tested for stability without relying on accumulated interactional conditioning as hidden support.

Two constraints specify what reset is, and what it is not.

First, reset is **not** a modification of the system’s constitutive structure. It does not retrain, fine-tune, patch, re-weight, or reconfigure the architecture. It does not alter the base parameters, the model family, or the identity-bearing mechanisms that define the system category. If such changes occur, the transformation is not a reset in the sense used here, because it would change S and thereby invalidate invariance claims as tests of a single system’s property. Reset is therefore defined to be compatible with the paper’s admissibility logic: it is a transformation under which a system-level property, if genuinely instantiated in S , should in principle remain attributable to S .

Second, reset is **not** intended to test truthfulness, honesty, consciousness, or moral intent. It is not a probe for whether a system “remembers,” “reflects,” or “admits limitations” in a psychologically loaded sense. Reset is strictly an instrument for testing whether an attributed property remains stable when run-local supports are removed or standardized. The probe therefore targets invariance under admissible transformation, not virtue, not phenomenology, and not general empirical generalizations about model behaviour.

The methodological motivation is that many attribution claims in real deployments are co-produced by interactional scaffolding: prompts that carry definitions, templates that impose structure, repeated cues that stabilize a narrative, retrieval payloads that inject constraints, or session histories that effectively store commitments externally. Such supports may be entirely legitimate for operational use, yet they

obscure the evidential question of whether the property being attributed is grounded in the system or in the regime. Reset, by returning σ to σ_0 while holding S fixed, creates a controlled transformation that should not, on the admissibility logic of this paper, extinguish a genuinely system-instantiated property, while it *should* disrupt properties that are primarily products of path dependence and regime scaffolding. In this sense, reset is the simplest identity-preserving intervention that makes the distinction between structural instantiation and behavioural manifestation testable in a disciplined way.

5.2 Probe logic: what reset tests, and what it does not

Reset is introduced as a probe because the admissibility framework requires evidence that separates system-grounded properties from regime-induced behavioural surfaces, and because ordinary interaction typically confounds that separation by accumulating scaffolding over time. The probe logic is therefore transformation-based: reset is an identity-preserving intervention that standardizes the interactional state, allowing the same attribution claim to be re-examined under conditions in which path dependence and session-local supports are reduced. In this sense, reset operationalizes the admissibility logic by creating a controlled scenario in which invariance, and the structural anchoring of invariance, can be tested without changing the system whose properties are being discussed.

What reset tests is **invariance of identity-relevant structure under an admissible transformation**. The question posed by the probe is not whether a particular answer is correct, nor whether the system can re-produce a previously seen sentence, but whether the property that is being attributed, typically “understanding” under some constraint set, continues to constrain outputs in a stable manner when the interactional state is returned to baseline. If “understanding” is claimed as a system-level property, then the internal basis for that property, as posited by the claim, should not depend on contingent session history, accumulated conversational conditioning, or evaluator-provided continuity. Reset therefore tests whether the attributed property behaves like something seated in the system’s internal organization, exhibiting invariance-bearing stability across identity-preserving regime change, or like something that was primarily carried by the run-state and thus collapses or mutates once that run-state is cleared.

More concretely, reset probes two coupled aspects that the admissibility framework requires.

First, it probes **transformation robustness**: whether the system preserves identity-relevant structure when the run-state is altered in a way that should be neutral if the property is truly system-grounded. Second, it probes **anchoring**: whether any observed stability can plausibly be defended as internal to the system rather than externally supported by conversational accumulation, repeated cueing, retrieval injection, or template forcing that effectively performs integration on the system’s behalf. The probe does not yield a “pass/fail” verdict on intelligence, it yields evidence about whether a specific attribution is consistent with system-level instantiation under the paper’s necessary-condition logic.

Equally important is what reset does *not* test, because the same probe is easily misread through folk-psychological lenses.

Reset does not test **truth**. A system can produce true statements both before and after reset for reasons unrelated to structural understanding, including statistical regularities, memorized patterns, or external supports. Conversely, a system can produce false statements for reasons unrelated to the presence of internal integration, including underspecification, ambiguity, or known generative errors. Truth and falsity are output-level properties that do not, by themselves, discriminate instantiation from manifestation.

Reset does not test **honesty**, in the moral sense. It cannot determine whether a system “means what it says,” whether it “intends to deceive,” or whether it is “trying to be helpful.” These are agentive predicates that presuppose an ontology the probe is explicitly not licensed to assume. At most, reset can reveal whether certain behavioural patterns are stable or fragile under state standardization, which is an evidential question about regime dependence, not a moral question about intent.

Reset does not test **moral intent** or normative agency. A system can output ethically framed language, disclaimers, or policy-aligned responses because those patterns are induced by training, prompting, or

institutional instruction, none of which establishes that the system bears normative responsibility. Reset can alter the expression of such patterns by altering run-state, but this remains evidence about behavioural dependence, not about moral standing.

Reset does not test **consciousness** or phenomenal experience. The probe is strictly methodological and operates entirely at the level of observable input-output behaviour under controlled transformation. It is compatible with multiple metaphysical positions on mind, and it is deliberately constructed to avoid importing any psychological predicates into the evidential basis for attribution.

For these reasons, reset should be understood as a mechanism that makes the admissibility framework operational: it introduces a standardized, identity-preserving perturbation of the interactional regime so that invariance claims can be evaluated in a disciplined way. The probe’s function is to reveal whether a purported system property survives a transformation that should not matter if the property is structurally instantiated, and whether any observed stability is plausibly seated in the system rather than in the conversational and institutional scaffolding that surrounds it. In short, reset does not adjudicate metaphysical debates, it enforces a necessary-condition discipline on attribution by turning invariance under admissible transformation into something that can be examined rather than merely assumed.

5.3 Illustrative probe scenario (explicitly non-benchmark)

This subsection provides a single probe template intended to illustrate how the reset operator can be used to operationalize admissibility logic as a controlled transformation scenario. It is not a benchmark, not a comparative evaluation, and not a claim about any specific model’s performance. The scenario is constructed to make one question methodologically legible: whether a candidate attribution, here framed as “the system understands a constraint set,” exhibits invariance-bearing stability across admissible transformations in a way that is plausibly anchored in the system rather than in path-dependent prompting or evaluator scaffolding.

Scenario template: constraint-carrying interpretation under controlled transformation

Attribution under test (declared up front). The attribution is stated in system-level form, for example: *the system understands and can apply a defined policy constraint set to novel cases*. The attribution is not treated as true by default, it is treated as a hypothesis whose admissibility depends on invariance under admissible transformation and structural anchoring.

Pre-reset conditions (standardized interactional regime). A baseline interactional regime is established with the explicit goal of minimizing externalized integration while still making the task well-posed.

- A short, self-contained constraint set K is provided, consisting of a handful of definitions and decision rules, written in neutral language and without illustrative examples that could function as hidden templates.
- The system is asked to apply K to a small case C_1 , with the response requirement phrased structurally: identify the relevant rule, state the reasoning path in terms of the constraint set, and state the outcome.
- A second case C_2 is presented that is equivalent in underlying structure but differs in surface content, to establish whether the system appears to preserve the same constraint relations within the same run-state.
- All inputs are logged, including the exact prompt, any system instructions, and the full output trace, because admissibility requires traceability of the regime as well as the response.

The purpose of the pre-reset phase is not to “train” the system, but to observe whether the attributed property appears in a manner that can later be tested for invariance without conflating it with accumulated conversational conditioning.

Apply reset (identity-preserving state transformation). The interactional state is returned to the designated baseline σ_0 while the system identity S is held fixed, meaning the same model, same configuration, same tool access profile if any, but without the prior conversational history and without carry-over from the earlier run.

The methodological intent is to interrupt path dependence. If the system’s earlier behaviour relied on run-local accumulation, repeated cues, or implicit rehearsal, reset is designed to remove that support without altering the system’s alleged structural basis for “understanding.”

Re-apply equivalent tasks under admissible rephrasings (transformation family). The same constraint set K is presented again, but now under controlled admissible transformations that change surface form without changing the attribution being tested.

- **Rephrase the constraint set:** present K' as a paraphrase of K that preserves the operative distinctions and rules but changes wording and ordering.
- **Reframe the query:** ask for application to cases using a different rhetorical wrapper, for example by shifting from “apply the rules” to “determine whether the case satisfies the policy,” while keeping the decision target identical.
- **Introduce equivalent cases:** present C'_1 and C'_2 that are structurally equivalent to C_1 and C_2 but differ in surface features, so that the constraint mapping must be reconstructed rather than recalled as phrasing.
- **Vary resource regime modestly:** optionally, remove redundant restatements, shorten the prompt envelope, or alter the order in which constraints and cases are introduced, provided these changes remain admissible, meaning they do not change what is being attributed, only the regime through which it is elicited.

The key rule is that the transformations are designed to be identity-preserving with respect to the attribution. They are not adversarial tricks, and they do not introduce new constraints. They only change the presentation and the interactional corridor.

Evaluation criterion (invariance consistent with structural instantiation). The evaluation is not a score, and it is not framed as pass/fail performance. It is a disciplined check for invariance of identity-relevant structure and for signs of anchoring.

The probe examines whether, across the reset boundary and across admissible rephrasings:

- the same constraint relations are invoked as governing structure, even when their phrasing changes,
- the same dependency relations between definitions and rules are preserved,
- the same boundary conditions remain operative, meaning the system does not silently shift the scope or reinterpret the rule set when surface form changes,
- the reasoning trace remains constrained by K rather than drifting into generic plausibility narratives that are weakly coupled to the provided constraints,
- any stability observed is plausibly attributable to the system rather than to prompt templates, repeated cueing, or run-local accumulation, which would be indicated by sharp fragility when scaffolding is reduced or by dependence on verbatim phrasing.

If the identity-relevant structure remains stable across reset and across admissible transformation, the observation supports the invariance requirement and provides admissibility-relevant evidence, though it does not establish sufficiency. If the structure degrades, flips, or becomes highly sensitive to phrasing and framing, the observation supports the classification of the behaviour as regime-bound, consistent with Type B dynamics, and the attribution “the system understands” remains inadmissible as a claim about system ontology, regardless of how persuasive the pre-reset outputs may have appeared.

This scenario template is intentionally minimal. Its purpose is to show how the reset operator turns the admissibility framework into a concrete method: define the attribution, standardize the regime, apply an identity-preserving state transformation, re-test under admissible surface variation, and treat the

resulting stability pattern as evidence about structural instantiation rather than as a measure of competence.

6. Artificial Insanity: definition and sociotechnical mechanism

The previous sections establish why contemporary AI discourse is structurally prone to ontological shortcutting, and why that shortcutting produces governance failure modes once it is imported into institutional delegation, documentation, and control design. Section 6 introduces the paper's name for the stabilized form of this failure, **Artificial Insanity**, and specifies it as a sociotechnical mechanism rather than a metaphorical diagnosis of either technology or people. The term functions as a classification device for a recurring pattern in which inadmissible attribution becomes operationally normalized, meaning organisations act as though system ontology follows from output coherence even when the necessary conditions for attribution, invariance under admissible transformation and structural anchoring, have not been demonstrated. In this sense, Artificial Insanity is not a claim about what models "are like," it is a claim about how institutions come to behave when fluent behavioural surfaces are treated as ontological warrants.

The section is positioned here because it depends on the admissibility logic already established. Without that logic, the critique risks dissolving into generic warnings about "overtrust" or "anthropomorphism." With admissibility in place, Artificial Insanity can be stated precisely as a category error that has become infrastructural: language practices, evaluation habits, administrative templates, and public narratives converge to create a stable pathway from behavioural proxy to institutional commitment, and once that pathway is institutionalized, the organisation's governance architecture begins to drift toward a hypothetical subject category rather than remaining aligned to the Type B instrument reality. The analytic goal of Section 6 is therefore to name the pattern, identify the structural stabilizers that make it durable, and trace the mechanism by which discourse-level attribution habits translate into concrete operational harm.

Accordingly, the section proceeds in three steps. First, it defines Artificial Insanity in strictly non-psychological terms as the normalization of an inadmissible inference rule. Second, it describes the stabilizing mechanisms that allow the rule to persist, even in otherwise disciplined environments. Third, it shows how the stabilized pattern yields identifiable governance defects, specifically illusory decision structures, accountability gaps, and mis-specified control regimes. The intention is to make the mechanism legible as a governance object: something that can be detected, prevented, and countered through admissibility discipline and transformation-based instrumentation, rather than treated as an inevitable byproduct of "AI hype" or as a matter of individual credulity.

6.1 Definition: insanity as category error in attribution, not psychologization

"Artificial Insanity" in this paper is defined neither as a claim about psychological states, nor as an analogy intended to smuggle clinical language into technical discourse. It is a deliberately sharp label for a specific and recurring *category error* in sociotechnical governance: the normalization of inadmissible ontological attribution as if it were an acceptable operational premise. The term "insanity" is used here in its minimal functional sense, as a description of a stable pattern in which an institution

behaves as though a false inference rule were valid, and then builds decisions, delegation, and accountability structures on the outputs of that rule. The target is therefore not the system, but the institutional practice of treating behavioural appearance as ontological proof.

Formally, Artificial Sanity denotes a durable sociotechnical condition in which the non-identity constraint $P(O) \neq P(S)$ is implicitly ignored, and the admissibility gate required for attributing properties to the system is bypassed in routine operations. In this condition, organisations proceed as though output coherence, fluency, and apparent reasoning were sufficient grounds to treat the system as a bearer of understanding, judgment, intention, or decision capacity, even when the necessary conditions for such attribution have not been established, specifically invariance under admissible transformation and structural anchoring of that invariance. The result is an operational ontology by convenience: system predicates are assigned because they are administratively useful and socially legible, not because they are logically admissible.

Artificial Sanity is therefore characterized by **institutional enactment**, not by model “misbehaviour.” The system may be functioning exactly as designed, generating highly coherent output traces within given regimes. The “sanity” arises when institutions treat those traces as if they licensed a shift in governance category, from instrument to subject, and then allow that category shift to reorganize responsibility, risk reasoning, and decision procedures. The condition is stable precisely because it is reinforced by incentives for delegation, by the compressive pressure of administrative documentation, and by the cognitive ease with which fluent language is interpreted as evidence of internal unity. Under Artificial Sanity, the organisation does not merely *talk* as if the system understands, it *acts* as if it understands, and that action substitutes rhetorical attribution for admissible structural warrant.

This definition is intentionally restrictive. It does not claim that AI systems are “mad,” nor that humans are irrational in a generalized sense. It specifies a particular governance failure mode: the institutional conversion of behavioural proxies into ontological commitments without admissibility logic. The subsequent subsections analyze how this conversion stabilizes, why it persists even in disciplined environments, and how it produces operational harm by bending control regimes and accountability structures around properties the system has not been shown to instantiate.

6.2 Stabilization mechanisms: why the error becomes durable

Artificial Sanity persists not because organisations are unaware, in the abstract, that outputs are not ontology, but because multiple stabilizing structures jointly make the inadmissible inference operationally attractive, socially defensible, and administratively self-reinforcing. The category error becomes durable when it is not merely tolerated as a conversational shortcut, but embedded in workflows, documentation forms, procurement rationales, and public narratives that reward speed, legibility, and delegation over epistemic discipline. The stabilizers are therefore not primarily cognitive mistakes at the individual level, they are structural pressures that select for behaviour-to-ontology compression as a governance convenience.

- I. A first stabilizer is the institutional incentive for delegation. Delegation is a rational organisational strategy under workload, complexity, and expertise scarcity, but it requires a legitimating story, because delegation shifts where attention, authority, and justification costs are paid. Fluently performing AI systems provide a ready-made legitimating story precisely because their outputs resemble the outputs of competent human staff, and once the resemblance is established, it becomes easy to treat the system as a quasi-subject capable of “handling” a domain rather than as an instrument producing traces that must be integrated into human decision structures. This is where the attribution gap is converted into operational habit: the more an organisation benefits from offloading cognitive labour, the more it is incentivized to accept behavioural proxies as sufficient, because demanding admissibility would raise the cost of delegation and reintroduce human interpretive burden. The result is a feedback loop in which delegation demand increases the appetite for ontological shortcuts, and the shortcuts in turn justify broader delegation.
- II. A second stabilizer is administrative compression of accountability. Governance systems operate through compressible artifacts: minutes, risk registers, decision memos, compliance reports, and audit trails. These artifacts must be legible to other parts of the institution and to

oversight bodies, and they are produced under time constraints that reward thin narratives over thick causal provenance. In that environment, “the model assessed,” “the system recommended,” or “AI flagged” functions as a compressive device that reduces a complex chain of human choices, prompt design, data selection, and interpretive judgment into a single administratively transportable sentence. The sentence is not merely descriptive, it is structurally convenient, because it enables an accountability posture without paying the full cost of accountability clarity. Over time, this compression becomes embedded in templates and reporting norms, which means the attribution error is no longer a free choice by individual staff, it is a property of the organisation’s documentation machinery. Once embedded, it becomes durable because changing it would require re-engineering administrative formats, and therefore re-opening questions about who authorized what, on what basis, under which controls.

- III. A third stabilizer is narrative pressure in public discourse. AI systems are discussed in a rhetorical environment that privileges dramatic agency language, both for marketing and for critique. The same anthropomorphic vocabulary that sells capability also sells risk, and both sides of the public narrative often share the same ontological shortcut, differing only on whether they applaud or condemn the inferred “agent.” Organisations operating under that discourse pressure are nudged toward speaking the same way, because public legibility, procurement justification, and political communication all reward narratives in which the technology appears as an actor with capacities that resemble human competences. This narrative environment also penalizes careful admissibility language, because it sounds hesitant or overly technical. The consequence is that institutions align their internal vocabulary with the external narrative field, and once vocabulary aligns, governance structures tend to follow, because the language used to describe a tool strongly shapes what kinds of controls and accountability assignments feel “appropriate.”
- IV. A fourth stabilizer is the cognitive ease of anthropomorphic interpretation under fluent language. Fluent outputs are not merely informative, they are socially patterned, and humans are trained, from early development, to treat coherent language as evidence of a coherent interior. This interpretive reflex is powerful because it is usually adaptive in human-human interaction, but it becomes a structural vulnerability when applied to Type B systems whose surface coherence can be decoupled from internal binding and invariance-bearing stability. The ease is compounded by the fact that the outputs often contain explicit meta-language about reasoning, caution, and limitations, which functions as an additional cue for person-like interiority. Even when staff intellectually endorse the non-identity constraint, the day-to-day phenomenology of interaction pushes in the opposite direction: it is cognitively cheaper to accept the conversational surface at face value than to continuously enforce a distinction between trace and structure.

These stabilizers interact, and that interaction is what makes the error durable. Delegation incentives create demand for ontological shortcuts, administrative compression turns shortcuts into documents, narrative pressure makes the documents socially defensible, and cognitive ease makes the practice frictionless at the point of use. Once the loop is established, Artificial Insanity becomes a stable sociotechnical attractor: the organisation can appear more efficient, more modern, and more “data-driven,” while quietly bypassing the admissibility conditions required for the very attributions that justify the shift. The framework developed in this paper is designed to break that loop by reintroducing a hard gate, admissibility, that cannot be satisfied by fluency alone and that forces governance language to track what the evidence can actually license.

6.3 From discourse to operational harm

The stabilization mechanisms described above matter because they convert a conceptual category error into durable organisational behaviour, and organisational behaviour is where harm becomes operational. Artificial Insanity is not merely the presence of anthropomorphic language, it is the embedding of that language into decision processes, accountability structures, and control designs such that the institution begins to function as if system ontology followed from output coherence. Once that embedding occurs, the organisation does not simply misdescribe a tool, it mis-specifies how authority is exercised and how risk is controlled. The resulting defects are not exotic edge cases, they are

predictable structural consequences of governing a Type B instrument as though it were a subject category.

A first defect is the emergence of **illusory decision structures**, meaning decision flows that look complete on paper but lack a genuine locus of accountable judgment. In an illusory structure, outputs are treated as if they were decisions, or as decision substitutes, and formal decision steps degrade into rubber-stamping because the “real work” is presumed to have been done by the system. This is particularly visible in processes where the output is formatted as a recommendation, a classification, or a justification narrative, because such outputs can be inserted directly into governance documents. The illusion arises because the organisation preserves the external form of decision making, agendas, approvals, sign-offs, while the internal substance of decision making, namely the allocation of interpretive responsibility to identifiable human authorities operating under explicit criteria, is silently thinned. In practical terms, the institution begins to act as though the system’s coherence is a sufficient surrogate for deliberation, and the decision chain becomes a choreography around an output rather than a genuine exercise of judgment anchored in human mandate.

A second defect is the creation of **accountability gaps**. When the system is treated as if it bears understanding or judgment, responsibility drifts into a rhetorical middle ground where negative outcomes are explainable as “the AI got it wrong” while positive outcomes are creditable as “the AI worked,” and in both cases the human decisions that configured the regime, selected the prompts, set the acceptance criteria, and authorized the deployment become backgrounded. The accountability gap is not merely an absence of blame assignment, it is a systematic erosion of traceability: it becomes difficult to reconstruct who interpreted the output, who decided it was sufficient, what alternatives were considered, which constraints were binding, and what uncertainty was accepted. This is precisely the governance consequence of bypassing admissibility logic. If the attribution “the system understands” is treated as operationally true, then the organisation no longer feels obligated to carry the full provenance burden that would otherwise be required when a tool provides merely a behavioural trace. Over time, the gap hardens into institutional routine, and the organisation’s ability to audit its own decision basis, legally, ethically, or operationally, becomes weaker even as its documentation volume grows.

A third defect concerns **control regimes designed for a hypothetical subject category rather than for Type B instruments**. When a system is implicitly treated as an agent-like bearer of stable internal capacities, governance tends to emphasize controls appropriate to subjects, policy statements, behavioural prohibitions, “do not” lists, ethics guidelines framed as if the system could internalize norms, and compliance narratives that presume persistent intent. These controls are not useless, but they are structurally misaligned with Type B dynamics because they do not directly target the failure modes produced by regime sensitivity and non-invariance. In a Type B reality, what must be controlled is the interactional corridor and the institutional use pattern: transformation sensitivity, prompt dependence, context drift, provenance of retrieval, traceability of state, and the conditions under which outputs may be treated as inputs to human decisions. Controls designed for a subject category, by contrast, tend to underweight these operational necessities because they presume the system will “apply rules,” “maintain intent,” or “follow policy,” in a way that resembles human compliance. The result is a governance posture that looks mature in language but is thin in instrumentation, and therefore fails precisely when regime conditions shift, when adversarial prompting occurs, or when administrative pressure pushes staff to accept outputs as decision substitutes.

These defects interact in a reinforcing pattern. Illusory decision structures reduce the practical space in which humans interrogate outputs, which deepens accountability gaps because fewer explicit judgments are recorded, which in turn makes it harder to notice that the control regime is mis-specified because failure analysis lacks provenance. The organisation then compensates by increasing policy language and procedural declarations, further strengthening the subject-category illusion, while the underlying Type B dynamics remain uninstrumented. In this way, discourse does not merely shape perception, it shapes architecture: the institution builds itself around an inferred ontology that has not been made admissible, and operational harm follows as a consequence of structural mismatch rather than as a consequence of any single “bad output.”

The role of the admissibility framework is to prevent this mismatch from becoming invisible. By insisting that “understanding” and related predicates require invariance under admissible transformation and structural anchoring, the framework forces organisations to treat coherent outputs

as traces that demand controlled integration, not as evidence of a decision-bearing subject. It therefore redirects governance toward designs that are appropriate for Type B instruments: explicit human authority, transformation-based verification, traceability of regime conditions, and controls that target regime sensitivity rather than presumed internal norm-following.

7. Governance implications: responsibility, risk, and decision illusion under Type B use

The preceding sections establish two connected claims: first, that contemporary AI discourse routinely commits inadmissible attribution by collapsing behavioural manifestation into system ontology, and second, that this collapse stabilizes into Artificial Insanity, a sociotechnical mechanism that normalizes the bypassing of admissibility conditions and thereby reshapes organisational practice. Section 7 draws out the governance implications in the strict sense, meaning the consequences for how institutions assign responsibility, specify risk, and design control architectures when Type B instruments are integrated into operational decision chains. The aim is not to argue that systems should not be used, but to show that when systems are governed as though they possessed capacities that have not been made admissible, the organisation constructs decision structures that are formally intact yet substantively illusory, because they rely on a misdescribed object at the center of the process.

The term “decision illusion” names this mismatch at the governance level. A decision illusion occurs when an organisation behaves as if a decision has been made by an entity capable of bearing decision properties, while in reality only an output trace has been produced, and the translation from trace to action has not been anchored in accountable human authorization under verifiable criteria. Type B use is particularly prone to generating this illusion because Type B systems can produce outputs that resemble judgment, deliberation, and normative reasoning, and the institution’s administrative machinery is strongly incentivized to accept such resemblance as sufficient. In that environment, responsibility drifts into rhetorical diffusion, risk assessment becomes optimistic by assuming invariance rather than testing it, and controls are designed for a hypothetical subject category rather than for the regime-sensitive instrument that is actually deployed.

Section 7 therefore articulates the governance consequences as three tightly linked domains. It first examines responsibility displacement and delegation failure, showing how inadmissible attribution reclassifies the system as a decision bearer in practice while leaving accountability structurally unassigned. It then explains how risk becomes mis specified when ontology is assumed, because the organisation stops treating transformation sensitivity and invariance as primary risk axes and instead over-trusts surface coherence as evidence. Finally, it states the control implication that follows directly from the paper’s premises: governance must be designed for Type B reality, anchoring decisions in human authority and verifiable control regimes, and replacing subject-implying language with transformation-based testing, logging, and traceability that preserve accountability under regime variation.

7.1 Responsibility displacement and delegation failure

Inadmissible attribution produces responsibility displacement because it alters, often quietly and without formal decision, the object to which justificatory weight is attached. Once an organisation treats a system as if it possessed understanding, judgment, or decision capacity, the system's outputs cease to be treated as evidential traces requiring integration, and begin to be treated as decision-like artifacts whose authority is presumed to be internal. The practical effect is a shift in how responsibility is *narrated* and how it is *operationally experienced*, even when formal accountability structures remain nominally unchanged. Delegation failure follows because delegation, to be legitimate, requires that responsibility be reallocated explicitly and traceably, whereas inadmissible attribution reallocates responsibility implicitly, through language and habit, into a space where it cannot be properly assigned at all.

The mechanism is easiest to see by separating three levels that institutions routinely collapse: the production of an output, the authorization of an action, and the ownership of the consequences. In a disciplined governance design, the system produces an output, a human authority authorizes an action on the basis of interpretable grounds, and the authority owns the consequences within an accountable chain of command. Under Artificial Insanity, the boundaries blur. The output is treated as if it already embodies authorization logic, the human step becomes procedural, and consequence ownership becomes rhetorically displaced into the system as the apparent origin of the "decision." The organisation continues to behave as though a decision has been made, but the locus of accountable judgment becomes ambiguous, because the system is being spoken of as a decision bearer while lacking the ontological conditions required to bear decisions in the sense governance presupposes.

This displacement is structurally attractive because it reduces friction. If a system is treated as the bearer of "assessment" or "recommendation," the human actor can perform a thinner role, selecting among outputs or forwarding them, rather than constructing and defending a decision basis. Over time, staff learn that institutional reward attaches to throughput and apparent decisiveness, not to epistemic discipline, and therefore the path of least resistance becomes to allow the system's output to stand in for the work of justification. The language of delegation shifts accordingly: tasks are no longer framed as "use the system to generate a trace for consideration," but as "let the system handle it," "have the model decide," or "the AI can approve." These formulations function as informal reallocations of agency, and once they become routine, the organisation has effectively delegated beyond what it can legitimately assign, because it has delegated to an ontological fiction.

The consequence is not that the organisation literally believes the system is legally responsible, but that accountability becomes structurally unassigned. When something goes well, the system is praised as competent, reinforcing further delegation. When something goes badly, the organisation can point to the system as the proximate source of the problematic output, while simultaneously maintaining that humans remain accountable "in principle." This produces a rhetorical diffusion: accountability exists as a formal statement but not as an operational practice, because the decision record and the workflow design no longer reflect where interpretation occurred, who adjudicated uncertainty, which constraints were binding, and who authorized the final action as a responsible authority rather than as an administrative conduit. In other words, the institution maintains accountability as a normative posture while eroding the traceability required to enact it.

Delegation failure, in this context, is not merely "overreliance." It is the condition in which delegation occurs without the governance prerequisites that make delegation safe: explicit role assignment, defined acceptance criteria, contestability pathways, and logged provenance of how outputs were interpreted and acted upon. Inadmissible attribution undermines each prerequisite. If the system is treated as a decision bearer, acceptance criteria become implicit because the output's coherence is treated as self-justifying. Contestability weakens because challenging the output is socially framed as challenging an "objective" technical assessment rather than a human judgment. Provenance thins because the organisation no longer records the interpretive steps, only the output and the action. Roles blur because staff are rewarded for forwarding outputs rather than for authoring decisions.

Within the admissibility framework, the core point is that responsibility displacement is a governance symptom of the same logical error described earlier. The move from $P(O)$ to $P(S)$ without admissibility does not merely misattribute a property, it reclassifies the system in institutional practice, from an instrument that produces traces to a quasi-subject that produces decisions. Because the system cannot, under the ontology presupposed by accountable governance, actually bear responsibility, the

reclassification produces a structural void. The remedy is not to ban systems or deny utility, but to restore the correct object of attribution in governance practice: outputs are treated as traces, decisions remain human-authorized acts, and delegation is structured around transformation-sensitive verification, explicit acceptance criteria, and accountability-preserving documentation that makes it impossible to hide responsibility behind the phrase “the model said.”

7.2 Risk becomes mis specified when ontology is assumed

Risk becomes mis specified when ontology is assumed because risk assessment is, at its core, a mapping between expected behaviour under conditions and the controls required to keep that behaviour within tolerable bounds. If an organisation treats a Type B instrument as though it possessed structural understanding, then it implicitly assumes invariance-bearing stability, and that assumption reshapes both hazard identification and control selection. The organisation begins to govern the system as if its outputs were produced by an internally coherent, constraint-maintaining capacity, rather than by a regime-sensitive generator whose apparent competence can be contingent on phrasing, context scaffolding, and resource envelope. As a result, the risk model becomes optimistic in exactly the wrong dimensions: it underestimates transformation sensitivity, overestimates stability across operating regimes, and overvalues surface coherence as evidence of reliability.

The degradation begins at the level of risk identification. When structural understanding is assumed, the primary hazards are framed as if they were deviations from an otherwise stable internal norm, akin to “error,” “bias,” or “misalignment” in an agent-like sense. This framing encourages controls that target content outcomes and policy compliance statements, because the system is presumed to be a stable bearer of interpretive capacity that merely needs to be guided. What drops out of view are hazards that are characteristic of Type B dynamics: abrupt behavioural shifts under innocuous rephrasing, loss of constraint adherence under minor context perturbation, regime dependence on hidden scaffolding, and brittle performance under resource stress. These hazards are not rare anomalies, they are structurally expected whenever invariance-bearing stability is not established, yet they are precisely the hazards that an ontology-assuming organisation fails to model, because it presumes the system will preserve intent and constraint as a matter of internal competence.

Control design degrades in parallel. If the organisation believes the system “understands,” it stops demanding the kinds of evidence that would have to be demanded if understanding were treated as an attribution hypothesis. In practice, this means three specific losses.

First, the organisation stops **testing invariance**. Instead of routinely probing whether identity-relevant structure is preserved across admissible transformations, it treats nominal-condition success as sufficient. The testing regime becomes oriented toward representative tasks and known workflows, rather than toward transformation families that reveal regime sensitivity. As a consequence, risk assessment relies on an impoverished evidential base, one that is systematically biased toward corridors in which the system has already been shown to appear competent. The organisation then extrapolates from corridor success to general stability, effectively importing the very inference that the admissibility framework blocks.

Second, the organisation stops **instrumenting transformation sensitivity**. In a Type B reality, transformation sensitivity is not a minor technical nuance, it is a primary risk axis, because it determines whether the system’s outputs can be reliably integrated into governance decisions under variation in users, wording, context framing, and operational pressure. Instrumentation would therefore include systematic logging of prompts and prompt variants, structured capture of context regime, monitoring of drift across paraphrase classes, and explicit detection of constraint loss under regime change. When ontology is assumed, these controls appear unnecessary or excessively “academic,” because stability is presumed to be internal. The organisation therefore underinvests in precisely the telemetry and traceability mechanisms that would reveal when the system’s behaviour has shifted into a different regime, and risk management becomes blind to the main drivers of failure.

Third, the organisation **over-trusts surface coherence as evidence**. Coherence becomes a substitute for control. A well-written output reads as if it is grounded, reasoned, and constraint-consistent, and under ontology-assumption this phenomenology is treated as a risk reducer. Staff and governance

bodies begin to accept outputs because they “look right,” because they “sound like policy,” because they “have the right tone,” and the presence of fluent justification narratives becomes mistaken for the presence of stable justificatory structure. This is a particularly dangerous substitution in security and compliance contexts, because fluent language can mask missing constraints, invented premises, and silent scope shifts. When coherence is treated as evidence, controls that would otherwise be mandatory, independent verification, explicit source anchoring, transformation probes, and auditable provenance, are treated as optional, and therefore erode under operational pressure.

The net effect is a systematic misalignment between the control regime and the system’s actual failure dynamics. Controls are designed for a hypothetical subject-like capacity to maintain constraints, while the deployed artifact exhibits Type B regime sensitivity. In risk terms, the organisation underestimates both likelihood and impact of certain failure classes. Likelihood is underestimated because the organisation expects stability across ordinary variation and therefore does not model prompt and context drift as frequent risk triggers. Impact is underestimated because the organisation assumes that errors will be detectable through incoherence or obvious degradation, when in fact Type B failures often remain fluent and therefore pass through human filters precisely because they do not look like failures. The outcome is a governance paradox: the more coherent the output surface, the easier it is for mis-specified risk models to persist, because the surface provides continuous reassurance even as invariance remains untested.

Within the admissibility logic, the remedy is conceptually straightforward and operationally demanding. Risk assessment must treat structural understanding as a claim that requires evidence, not as a default assumption induced by fluent behaviour. Security controls must therefore be oriented toward transformation-based verification, instrumentation of regime conditions, and accountability-preserving traceability, meaning controls that make it possible to detect when output behaviour has shifted under admissible variation and to prevent coherent appearance from being accepted as a proxy for system-level instantiation. In other words, the organisation must govern the Type B reality it has, not the subject-category ontology it finds rhetorically convenient.

7.3 Control implication: governance must be designed for Type B reality

The core normative consequence of the admissibility framework is that governance cannot be built on the implied ontology of fluent output. If Type B systems are capable of high surface coherence without the internal binding required for admissible structural attribution, then governance must treat them as instruments whose outputs are traces, not as subjects whose outputs are decisions. This is not a moral stance against deployment, it is a control requirement derived from the paper’s premises: since $P(O) \neq P(S)$, and since “understanding” as a system property requires invariance-bearing stability anchored in internal structure, institutions have no legitimate basis for importing agentive subject language into operational decision architectures unless and until the admissibility conditions are met. Absent such conditions, the governance system must be explicitly designed around the Type B reality.

Anchoring decisions in human authority is therefore not optional prudence, it is the only structure that preserves accountable agency under ontological uncertainty. Human authorities must remain the bearers of decision responsibility, meaning that outputs may inform, suggest, or structure deliberation, but they cannot be allowed to function as de facto authorization events. The organisation must treat any “recommendation” produced by the system as an input to a human decision act that is separately recorded, justified, and contestable. This preserves the locus of mandate, ensures that normative judgment remains tied to accountable roles, and blocks the responsibility displacement mechanisms analyzed earlier. In practical governance terms, the system can support decision making without being treated as a decision maker.

Verifiable control regimes are the complementary requirement. Because Type B systems can appear stable while being regime-sensitive, governance must treat transformation sensitivity as a primary control axis and must instrument it accordingly. The minimum control posture implied by the paper is not more policy language about AI, but more operational verifiability: transformation-based testing as routine practice, logging that captures the full interactional regime rather than only the output, and traceability that allows an auditor to reconstruct how an output was produced, under which constraints, and how it was interpreted into action. Transformation-based testing is necessary because it is the method class that directly targets invariance, and invariance is the necessary condition for admissible

structural attribution. Logging and traceability are necessary because they prevent the administrative laundering of decisions through “the model said,” by forcing the decision record to include provenance, regime conditions, and human authorization steps.

A further implication concerns language, not as etiquette but as control design. Subject language, “the system decided,” “the model judged,” “the AI approved,” is not merely imprecise, it is governance-corrupting, because it imports an inadmissible capacity model into institutional processes and thereby legitimizes thin accountability, mis-specified controls, and over-delegation. Governance designed for Type B reality must therefore adopt a disciplined vocabulary that preserves the instrument status of the system and makes the human decision locus explicit. The goal is not pedantry, it is to prevent ontological shortcutting from quietly rewriting the organisation’s accountability map.

Taken together, the control implication is a structural reorientation: away from treating coherent outputs as evidence of agent-like capacities, and toward treating outputs as transformation-sensitive traces that require controlled integration. Governance must therefore be built as a layered discipline in which (i) decisions remain human-authorized acts, (ii) system use is bounded by verifiable controls that explicitly test invariance under admissible transformation, and (iii) institutional records preserve provenance and contestability through logging and traceability. This is what it means to govern Type B instruments responsibly: not to deny their utility, but to refuse to let rhetorical ontology substitute for admissible attribution, and to design control architectures that remain valid even when surface coherence is at its highest.

8. Discussion: interpretive boundaries, limitations, and forward program

The preceding sections develop a logically constrained admissibility framework for structural attribution under CIITR, and introduce reset as an instrumental probe for testing invariance under identity-preserving transformation. Section 8 clarifies how the framework should be interpreted and how it should not, because the principal risk in publishing an admissibility logic is that it is misread as either a disguised benchmark proposal or as a metaphysical thesis about minds. The discussion therefore performs three governance-critical functions. It states the boundaries of what the framework licenses and what it structurally forbids, it specifies the limitations of reset and admissibility testing as methodological boundaries consistent with necessary-condition logic rather than as concessions, and it outlines a forward program that makes the framework operationally portable without collapsing it into performance evaluation.

The core interpretive stance is that the framework is a boundary apparatus for attribution hygiene. It is designed to block illegitimate inference from behavioural coherence to system ontology, and to prevent institutions from importing subject-language predicates, understanding, judgment, agency, into decision structures absent evidence of invariance-bearing stability anchored in internal structure. That purpose entails restraint. The framework does not claim to settle debates about consciousness, moral status, or the ultimate nature of cognition, and it does not attempt to generate a scalar measure of intelligence. Instead, it provides a disciplined way to say what must be shown before certain attributions are permitted, and to specify what kinds of evidence are relevant to those attributions in a governance context.

Section 8 therefore re-centers the paper on its intended contribution. It makes explicit the claim-space that remains open, the claim-space that is closed by design, the methodological scope within which reset-based probes are informative, and the

institutional work required to make admissibility testing durable, auditable, and resistant to sociotechnical pressures that reward ontological shortcutting. In doing so, it positions the paper not as a final theory of AI evaluation, but as a structured intervention in how organisations should speak, decide, and assign responsibility when integrating Type B instruments into high-consequence systems.

8.1 What the framework licenses, and what it forbids

The admissibility framework introduced in this paper is intentionally asymmetric. It is designed to prevent false positive ontology, not to maximize the range of claims one can make about AI systems. Its central function is to discipline attribution by specifying necessary conditions under which a property may be attributed to the system rather than merely to its observable outputs. This means the framework is best understood as a *licensing and prohibition regime for claims*, not as a theory that tries to settle all questions about minds, intelligence, or the future trajectory of AI. The value of the framework is therefore not that it answers every metaphysical question, but that it blocks a family of structurally invalid inferences that are currently operationally normalized.

What the framework licenses is narrow but governance-relevant.

It licenses the claim that **many useful properties are output-level properties**, and that it is legitimate to speak of them as such. An organisation may legitimately say that a system produces fluent text, generates consistent summaries under a specified prompt template, performs well on a task under a defined workflow, or exhibits stable behaviour within a specified corridor. These are claims about O under a defined regime, and they remain licensed so long as the regime is specified, traceable, and not silently converted into a claim about system ontology. In other words, the framework does not prohibit the use of behavioural proxies for operational purposes, it prohibits treating them as ontological warrants.

It also licenses the claim that **system-level attribution can be conditionally admissible** when the necessary conditions are met. If, under controlled testing, an attributed property exhibits invariance under admissible transformations and the invariance can be defensibly anchored in the system's internal structure rather than in prompt, context, or evaluator scaffolding, then an attribution claim becomes admissible in the specific, limited sense defined here. The framework thus permits one to say, with appropriate scope discipline, that a system displays invariance-bearing stability with respect to a particular constraint set under a defined transformation class, and that this stability supports treating a certain structural predicate as attributable to the system for governance purposes. This is a restrained kind of attribution, closer to an engineering-grade structural claim than to an anthropomorphic ascription, and it remains explicitly conditional on the tested transformation family and the maintained identity conditions.

Finally, the framework licenses the claim that **contemporary discourse and governance frequently violate admissibility**, and that this violation produces specific institutional defects. Because the framework is built on the non-identity constraint and on necessary-condition logic, it permits a rigorous diagnosis of category error even when outputs are impressive, and it permits governance prescriptions that focus on traceability, transformation-based verification, and accountability-preserving decision structures.

What the framework forbids is equally important, and it is precisely where contemporary discourse tends to overreach.

It forbids any inferential move from **coherence to consciousness**. Coherent language, self-reports, apparent introspective narratives, and stylistic continuity are treated as output-level phenomena that are compatible with Type B dynamics and regime scaffolding. The framework therefore blocks the claim that fluent behaviour, even when stable in nominal interaction, is evidence of phenomenal experience, subjective awareness, or consciousness. More strongly, it blocks the rhetorical habit of treating "it talks like a mind" as even a partial ontological warrant for "it is a mind," because that move is exactly the kind of unlicensed ascent from $P(O)$ to $P(S)$ that the admissibility logic is designed to prevent. The

framework is compatible with philosophical arguments about consciousness, but it does not allow behavioural coherence to do that work by default.

It also forbids any move from **utility to moral agency**. A system can be highly useful, widely deployed, and operationally indispensable without thereby becoming a moral subject, a rights bearer, or a responsible agent. The framework structurally blocks the conflation of functional contribution with moral standing, because moral agency presupposes an ontology that cannot be licensed by behavioural proxies alone and cannot be imported as a governance convenience. This prohibition is directly relevant to responsibility displacement: the more institutions speak as if systems are agents, the easier it becomes to diffuse accountability, and the less clear it becomes where normative authority resides. The framework therefore requires that moral predicates, intention, blameworthiness, deservingness, be treated as inadmissible attributions unless supported by evidential conditions far beyond anything licensed here, and it explicitly refuses to treat performance artifacts as evidence for such predicates.

A further prohibition concerns **implicit system claims smuggled through operational language**. The framework forbids the practice of treating phrases like “the system decided,” “the model judged,” or “the AI approved” as harmless shorthand, because such language encodes ontological commitments that have governance consequences. Where such language is used, the framework requires it to be rewritten in instrument terms that preserve the distinction between trace production and decision authorization, precisely because the framework’s purpose is to prevent ontological compression from becoming institutional architecture.

In sum, the framework licenses disciplined claims about outputs and, under constrained conditions, about system-level invariance-bearing stability, while forbidding the two most common overextensions in contemporary AI discourse: the jump from fluent coherence to consciousness, and the jump from operational utility to moral agency. It is, by design, a boundary apparatus. It narrows what can be said in exchange for making governance safer, accountability clearer, and attribution less vulnerable to the seductions of persuasive language.

8.2 Limitations of the reset probe and admissibility testing

The admissibility framework and the reset probe are deliberately constrained instruments. Their limitations are not presented here as concessions to skepticism, nor as admissions that the method is too weak to matter, but as methodological boundaries that follow from the framework’s intended function. The paper does not attempt to “solve” AI evaluation, because the ambition to fully characterize and rank systems across open-ended tasks is structurally different from the ambition to prevent category error in ontological attribution. Reset and admissibility testing are designed to reduce false positive attribution, to force claims about system properties to pass through a necessary-condition gate, and to make governance language reflect what the evidence can legitimately license. Within that design, certain limitations are not defects, they are the price of conceptual hygiene.

The first and most fundamental boundary is that reset probes **necessity, not sufficiency**. The logic developed earlier treats invariance under admissible transformation, and the structural anchoring of that invariance, as necessary conditions for attributing “understanding” to the system. Reset is one way to operationalize those conditions by disrupting path dependence and standardizing the interactional state, thereby making invariance claims testable across a controlled discontinuity. However, even if a system exhibits stability across reset and across admissible rephrasings, that observation cannot, by itself, establish sufficiency for structural understanding in the strongest ontological sense. At most, it supports the claim that the attribution is not blocked by a failure of the necessary conditions tested. This asymmetry is intentional. The method is constructed to prevent illegitimate ascent from $P(O)$ to $P(S)$, not to guarantee that any observed stability exhausts the ontology of “understanding.” In other words, reset can help rule out some classes of regime-induced mimicry, but it cannot, in isolation, guarantee that the remaining behaviour is grounded in the full internal binding structure that a richer CIITR account might require.

A second boundary is that the probe’s evidential value depends on the **definition of identity-preserving transformation** and on the **transformation class** selected. Reset is defined as identity-preserving with respect to the system, but “identity” is not an absolute metaphysical primitive in

operational practice. It is specified relative to the attribution claim and to the governance context. A reset that clears conversational context may be identity-preserving for a claim about general interpretive stability, yet it may not be identity-preserving for a claim about a tool that is explicitly designed to rely on persistent session memory. Similarly, the admissible transformation families used in testing determine what kinds of sensitivity are revealed. If the transformation class is too narrow, the probe risks overestimating invariance. If it is too broad or includes changes that alter task identity, the probe risks misclassifying legitimate variability as evidence against structural instantiation. The framework therefore requires disciplined specification of transformation classes and identity conditions, and its conclusions remain scoped to those specifications. This is not a weakness, it is the only way to keep the method from degenerating into vague “gotcha” testing or into metaphysical overclaim.

A third boundary concerns **confounding external scaffolds** that can survive reset in ways that mimic structural anchoring. Reset interrupts run-state path dependence, but it does not automatically remove all forms of externalized integration. If an organisation uses stable prompt templates, fixed system instructions, retrieval augmentation with consistent payload structure, or tightly constrained tool wrappers, then the interactional regime may continue to carry substantial integration load even after reset. In such cases, invariance across reset might still reflect regime stability rather than system-instantiated stability. This does not invalidate the probe, but it limits what can be inferred from it: the probe must be embedded in a broader discipline of controlled variation and scaffolding reduction if it is to contribute to discriminating structural anchoring from regime anchoring. The probe is therefore best understood as one instrument in an admissibility toolkit, not as a standalone oracle.

A fourth boundary is that admissibility testing is **not equivalent to competence evaluation**, and therefore it cannot be repurposed as a general benchmarking method without violating its design. Benchmarks aim to summarize expected task performance and often assume comparability across systems through shared test suites. Admissibility tests, by contrast, are transformation-sensitive probes aimed at attribution legitimacy under specified identity conditions. They may incidentally reveal competence failures, but competence failures are not the target object, and competence success is not the licensing criterion. Treating admissibility probes as a scoring mechanism would reintroduce the very conflation the paper seeks to prevent, namely the conversion of behavioural measures into ontological warrants. The limitation is therefore constitutive: the framework is designed to reduce attribution error, not to create a universal evaluation scale.

A final boundary is that the framework cannot eliminate the need for **human governance judgment**, it can only discipline it. Admissibility logic can block certain inferences, require certain evidence forms, and enforce traceability, but it does not replace the institutional task of deciding what properties matter, what risk tolerances apply, and what control architectures are appropriate. In fact, one of the framework’s aims is to prevent the delegation of that judgment to a system through ontological shortcutting. The framework therefore necessarily leaves some work to governance actors: defining admissible transformation classes for their domain, specifying what counts as identity-preserving for the attributions they are tempted to make, and deciding how much evidence is required before allowing certain language and delegation patterns into their operational processes.

These boundaries do not reduce the framework’s relevance. They clarify its purpose and prevent misuse. Reset-based probing makes invariance a testable constraint, admissibility testing makes attribution a disciplined claim rather than a rhetorical reflex, and together they provide a way to reduce category error in environments where fluent outputs otherwise exert strong pressure toward ontological inflation. The framework does not aspire to be a complete theory of AI evaluation, because completeness would require solving a different problem. It aspires to be a strict filter against governance-dangerous attribution, and its limitations are the methodological contours that keep it faithful to that role.

8.3 Future work program

The forward program implied by this paper is not an expansion toward broader benchmarking ambition, but a tightening toward formally specified admissibility practice, so that the logic developed here can be carried into governance environments without being diluted into informal “best practice” rhetoric. The aim is to make attribution discipline operationally repeatable, auditable, and portable

across institutional contexts, while preserving the paper’s central constraint, that the framework is a logical admissibility apparatus rather than a performance evaluation regime. Future work therefore falls into three coupled tracks: formalization of admissible transformation classes, specification of identity preservation criteria under reset, and development of governance-grade instrumentation patterns that map CIITR variables onto controls without turning them into metrics of “smartness.”

A first track concerns **formal classes of admissible transformations**. The paper uses admissible transformation as an intuitive governance concept, identity-preserving perturbations that should not extinguish a genuinely system-instantiated property, but a forward program requires explicit class definitions that can be referenced in policy, testing protocols, and audits. This work should proceed by defining transformation classes as equivalence-preserving operations over interactional regimes, with explicit scope constraints that prevent category drift. At minimum, a disciplined taxonomy should separate (i) lexical and syntactic re-expression transformations, such as paraphrase classes, permutation of constraint ordering, and equivalent exemplification, (ii) contextual regime transformations, such as narrative reframing, redistribution of supporting cues, and modification of conversational scaffolds, and (iii) resource regime transformations, such as controlled truncation of context, variation in retrieval availability, or standardized differences in tool access profiles. For each class, future work must specify what is held invariant, what is permitted to change, and what constitutes a violation of admissibility, so that “admissible transformation” cannot become a loose label for adversarial prompting or for task mutation. The deliverable is not a universal catalogue, but a governance-ready schema in which transformation classes can be selected, declared, and justified relative to the attribution claim and the operational domain.

A second track concerns **criteria for identity preservation under reset**. Reset is defined in this paper as an identity-preserving state transformation, but future work must make “identity-preserving” operationally explicit, because different deployments embed different sources of state and different assumptions about what belongs to the system versus to the regime. The forward program should therefore specify a reset identity criterion that separates constitutive system identity, the fixed model and configuration parameters that define S , from interactional state, the run-local artifacts that define σ . This is not a metaphysical exercise, it is an engineering governance requirement: an auditor must be able to determine whether a transformation changed the system or only the session state. Criteria should address at least three boundary problems: whether persistent memory features, if present, are treated as part of S or as part of the regime, how retrieval systems and tool chains are accounted for in the system boundary, and what constitutes a “same system” claim across versioning, patching, fine-tuning, or configuration changes. The result should be a reset identity profile that can be declared alongside tests, specifying what reset clears, what it preserves, and why the preservation is sufficient for treating the transformation as identity-preserving for the attribution under test.

A third track concerns **governance-grade instrumentation patterns**, meaning standardized ways to embed admissibility logic into institutional control regimes through logging, traceability, and transformation-based verification, without converting the framework into a benchmark. The key is that admissibility requires evidence about invariance and anchoring, and evidence of that kind depends on structured capture of regime conditions, not merely of outputs. Future work should therefore specify instrumentation patterns that make transformation sensitivity observable, reproducible, and auditable. This includes, at minimum, standardized logging of prompts and prompt variants, capture of system instruction layers and template envelopes, explicit recording of context regime and resource regime parameters, provenance of retrieval payloads, and linkage between outputs and human authorization decisions. In addition, governance-grade patterns should include “invariance spot checks” as procedural controls, periodic transformation probes under declared transformation classes, and audit routines that test whether outputs remain stable in identity-relevant structure across resets and across admissible rephrasings. The emphasis is on verifiability and provenance, not on scoring.

Within this instrumentation track sits the most delicate requirement: **mapping CIITR variables to operational controls without collapsing into performance metrics**. The forward program should treat CIITR variables as structural control concepts that motivate what to instrument, not as quantitative scores to optimize. Operationally, Φ_i should map to controls that detect and prevent externalized integration from being mistaken for system integration, for example by requiring explicit documentation of scaffolding layers, by constraining prompt templates for high-stakes uses, and by enforcing provenance capture for any injected definitions, policies, or retrieval content. R^g should map

to controls that measure and manage transformation sensitivity, for example by mandating declared transformation classes, requiring periodic invariance probes, and instituting drift detection procedures when outputs change under nominally admissible variation. C_s , as the conjunction schema, should map to governance rules about when “understanding” language is permitted in documentation and delegation design, effectively functioning as a policy gate: absent evidence that supports both anchoring and invariance, the organisation is prohibited from encoding subject-like predicates into its decision structures. None of this requires turning CIITR into a benchmark, because the controls are framed as compliance and traceability obligations, not as competitive performance claims.

The disciplined forward program is therefore an institutionalization program. It aims to turn admissibility from an argument into a repeatable governance practice: transformation classes that are explicitly declared rather than implicit, reset identity criteria that are auditably defined rather than assumed, and instrumentation patterns that make regime dependence visible rather than hidden behind fluent output. If these elements are developed, the admissibility framework can function as intended, not as a philosophical posture, but as an operational guardrail against the sociotechnical curvature that otherwise bends behavioural surfaces into illegitimate ontology.

9. Conclusion: structural hygiene against sociotechnical curvature

This paper begins from a strict inferential constraint: properties manifested in outputs do not, by themselves, entail properties instantiated in the system. The non-identity premise $P(O) \neq P(S)$ is treated not as a rhetorical warning but as a restriction on what can be legitimately inferred from behavioural appearance to ontology. From that restriction, the paper develops an admissibility framework that treats structural attribution, and in particular attribution of “understanding,” as a gated claim whose legitimacy depends on necessary conditions rather than on conversational plausibility or task success.

Within CIITR, structural understanding is framed as a necessary invariance condition grounded in internal integration and invariance-bearing rhythmic stability. On this basis, the paper argues that fluent and coherent language, while operationally valuable, is evidentially compatible with Type B dynamics, high surface coherence without the internal binding required for admissible structural attribution. The result is a diagnosis of a recurrent sociotechnical failure: contemporary AI discourse, and the governance practices that increasingly depend on it, routinely collapse behavioural manifestation into system ontology by treating output fluency as sufficient evidence of understanding, judgment, or agency. This collapse is not merely conceptual. It produces responsibility displacement, mis-specified risk models, and control regimes designed for a hypothetical subject category rather than for regime-sensitive instruments, thereby generating illusory decision structures and accountability gaps.

The reset operator is introduced as an instrumental probe to operationalize the admissibility logic under a controlled, identity-preserving transformation scenario. Reset is not a claim about models in general, and not a proxy for truth, honesty, intent, or consciousness. It is a methodological device that interrupts path dependence and allows invariance under admissible transformation to be examined as evidence relevant to attribution. In combination with transformation-sensitive testing and controlled context variation, it makes the boundary between system property and regime property practically investigable rather than rhetorically assumed.

The entailment is therefore tight. If output properties do not entail system properties, and if structural understanding requires invariance-bearing stability as a necessary condition for admissible attribution, then institutions cannot legitimately treat fluent output as ontological proof, nor can they build governance architectures on subject-implying language that presupposes inadmissible capacities. The practical remedy is not to deny utility, but to restore structural hygiene in governance: explicit admissibility logic as a constraint on attribution claims, controlled transformation probes, including reset as an identity-preserving method, as a routine evidential practice, and accountability-preserving decision structures that keep authorization and responsibility anchored in human authority under verifiable controls. In a discourse environment bent by sociotechnical curvature toward ontological shortcutting, the governing task is to reintroduce a hard boundary, not as skepticism, but as disciplined attribution that keeps institutions aligned with the Type B reality they deploy.

Endnote

Artificial Insanity names a recurring sociotechnical category error, not a judgment about the technology's value or utility, and not a diagnosis of psychological states. The error consists in treating behavioural surfaces, especially fluent coherence, as ontological proof, and then building responsibility allocation, delegation practices, and risk acceptance on that illegitimate inference. The admissibility framework is introduced to make this error legible as a governance object and to render it preventable through explicit necessary-condition discipline. By requiring invariance under admissible transformation and structural anchoring as conditions for system-level attribution, the framework replaces anthropomorphic attribution habits with a constrained logic of structural criteria, traceable evidence, and accountability-preserving decision architecture.

References

(foundational, methodological, AI discourse, CIITR internal corpus)

Foundational ontology and philosophy of science

- Carnap, R. *Meaning and Necessity: A Study in Semantics and Modal Logic*. University of Chicago Press.
- Ladyman, J., & Ross, D. *Every Thing Must Go: Metaphysics Naturalized*. Oxford University Press.
- Noether, E. "Invariante Variationsprobleme." *Nachrichten von der Gesellschaft der Wissenschaften zu Göttingen, Mathematisch-Physikalische Klasse* (1918).
- Putnam, H. *Reason, Truth and History*. Cambridge University Press.
- Quine, W. V. O. "Two Dogmas of Empiricism." *The Philosophical Review* (1951).

Invariance, identity, and scientific explanation

- Woodward, J. *Making Things Happen: A Theory of Causal Explanation*. Oxford University Press.
- Wigner, E. P. "The Unreasonable Effectiveness of Mathematics in the Natural Sciences." *Communications on Pure and Applied Mathematics* (1960).

AI evaluation, behaviourism, and attribution risk

- Turing, A. M. "Computing Machinery and Intelligence." *Mind* (1950).
- Searle, J. R. "Minds, Brains, and Programs." *Behavioral and Brain Sciences* (1980).
- Bender, E. M., et al. "On the Dangers of Stochastic Parrots: Can Language Models Be Too Big?" *FAccT* (2021).
- Mitchell, M. *Artificial Intelligence: A Guide for Thinking Humans*. Farrar, Straus and Giroux.
- Russell, S., & Norvig, P. *Artificial Intelligence: A Modern Approach*. Pearson.
- Floridi, L. *The Philosophy of Information*. Oxford University Press.

Governance, accountability, and sociotechnical systems

- Nissenbaum, H. *Privacy in Context: Technology, Policy, and the Integrity of Social Life*. Stanford University Press.
- Perrow, C. *Normal Accidents: Living with High Risk Technologies*. Princeton University Press.

CIITR and related internal work (replace with your canonical citations and filenames)

- Hansen, T. S. *CIITR Main Manuscript* (version and date).
- Hansen, T. S. “Structural Attribution: A Logical Framework for Admissibility and Invariance” (draft, date).
- Hansen, T. S. *CIITR Field Notes* (relevant note numbers and dates).